

# Lecture 16: Detection + Segmentation

# Reminder: A4

A4 due **Wednesday, November 13, 11:59pm**

A4 covers:

- PyTorch autograd
- Residual networks
- Recurrent neural networks
- Attention
- Feature visualization
- Style transfer
- Adversarial examples

# Last Time: Computer Vision Tasks


## Classification



**CAT**

No spatial extent


## Semantic Segmentation



**GRASS, CAT, TREE, SKY**

No objects, just pixels

## Object Detection



**DOG, DOG, CAT**

Multiple Objects

## Instance Segmentation



**DOG, DOG, CAT**

[This image is CC0 public domain](#)

# Last Time: Object Detection


Classification



CAT

No spatial extent


Semantic Segmentation



GRASS, CAT, TREE,  
SKY

No objects, just pixels

Object Detection



DOG, DOG, CAT

Multiple Objects

Instance Segmentation



DOG, DOG, CAT

# Object Detection: Impact of Deep Learning


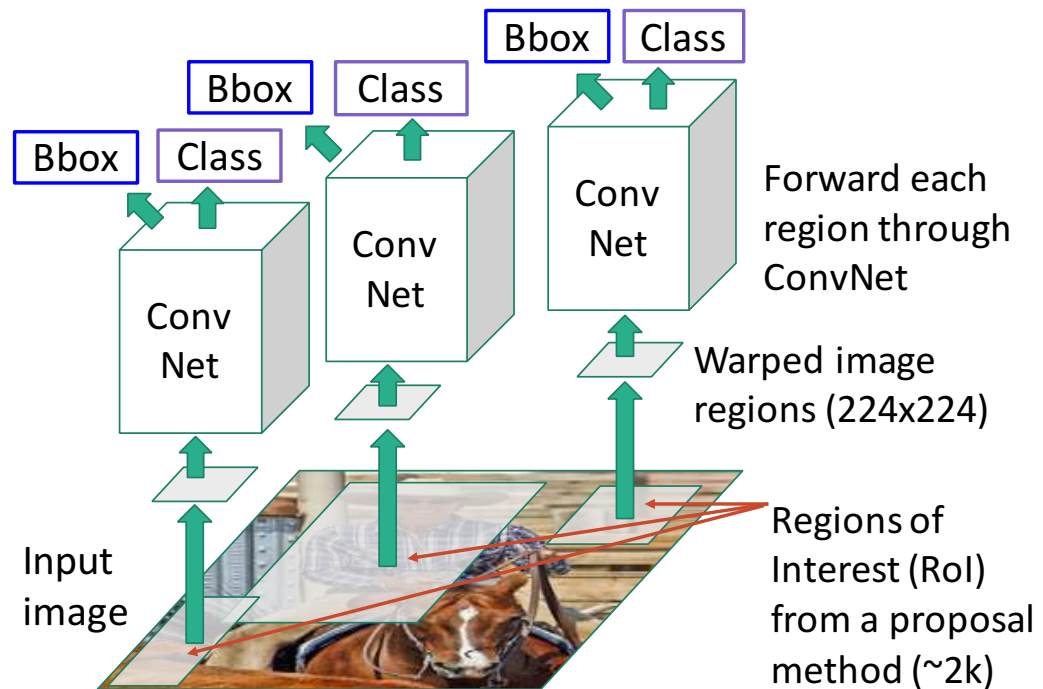



Figure copyright Ross Girshick, 2015.  
Reproduced with permission.

# Last Time: Object Detection Methods

**“Slow” R-CNN:** Run CNN independently for each region



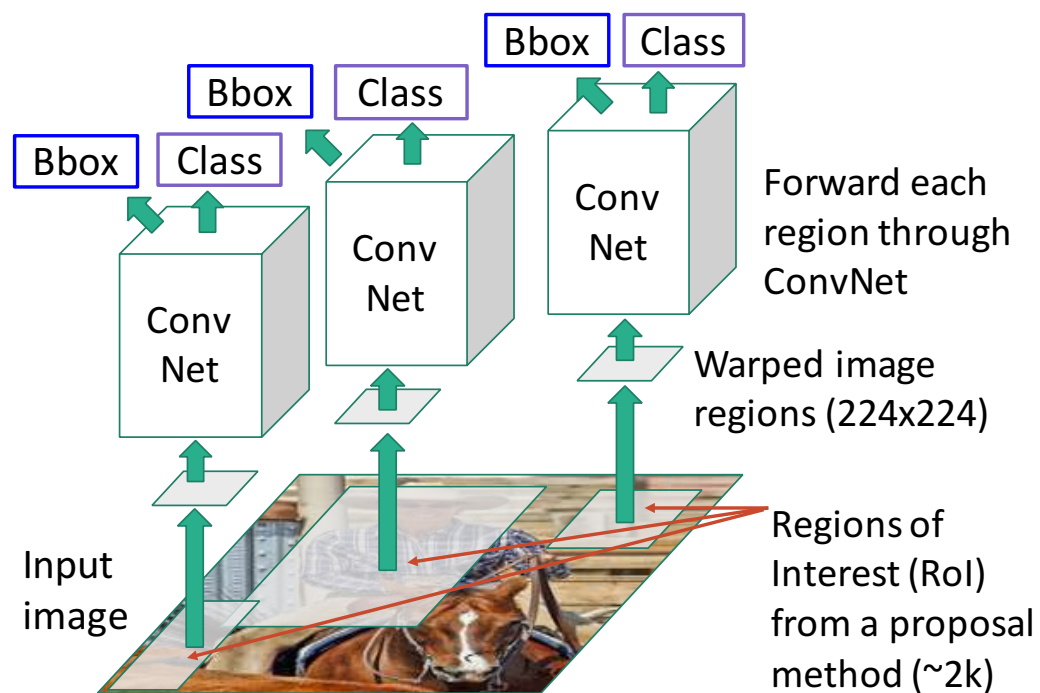
**Fast R-CNN:** Apply differentiable cropping to shared image features



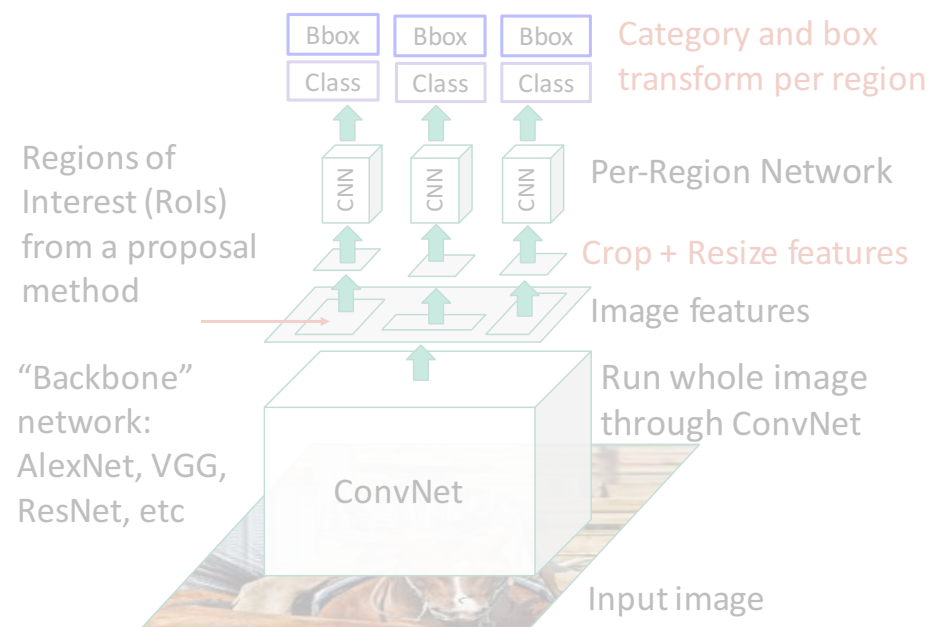


# Recap: Slow R-CNN Training

**“Slow” R-CNN:** Run CNN independently for each region




**Fast R-CNN:** Apply differentiable cropping to shared image features



# "Slow" R-CNN Training

Input Image




Ground-Truth boxes

[This image](#) is [CC0 public domain](#)



# "Slow" R-CNN Training

Input Image




Ground-Truth boxes

Region Proposals

[This image is CC0 public domain](#)

# "Slow" R-CNN Training

Input Image



GT Boxes

Positive

Neutral


Negative

Categorize each region proposal as positive, negative, or neutral based on overlap with ground-truth boxes

[This image is CC0 public domain](#)

# "Slow" R-CNN Training

Input Image





GT Boxes

Positive

Neutral

Negative




Crop pixels from  
each positive and  
negative proposal,  
resize to 224 x 224

This image is [CC0 public domain](#)



# "Slow" R-CNN Training

Input Image





GT Boxes

Positive


Neutral

Negative



Run each region through CNN. For positive boxes predict class and box offset; for negative boxes just predict background class




Class target: Dog  
Box target: →



Class target: Cat  
Box target: →



Class target: Dog  
Box target: →




Class target: Background  
Box target: None

[This image is CC0 public domain](#)

# Fast R-CNN Training

Crop features for each region, use them to predict class and box targets per region

Input Image



Backbone CNN


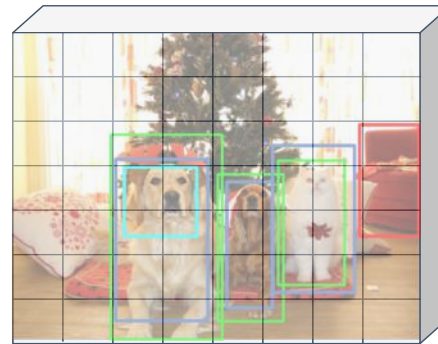


Image Features





GT Boxes


Positive

Neutral

Negative





Class target: Dog  
Box target: →



Class target: Cat  
Box target: →



Class target: Dog  
Box target: →



Class target: Background  
Box target: None

[This image is CC0 public domain](#)

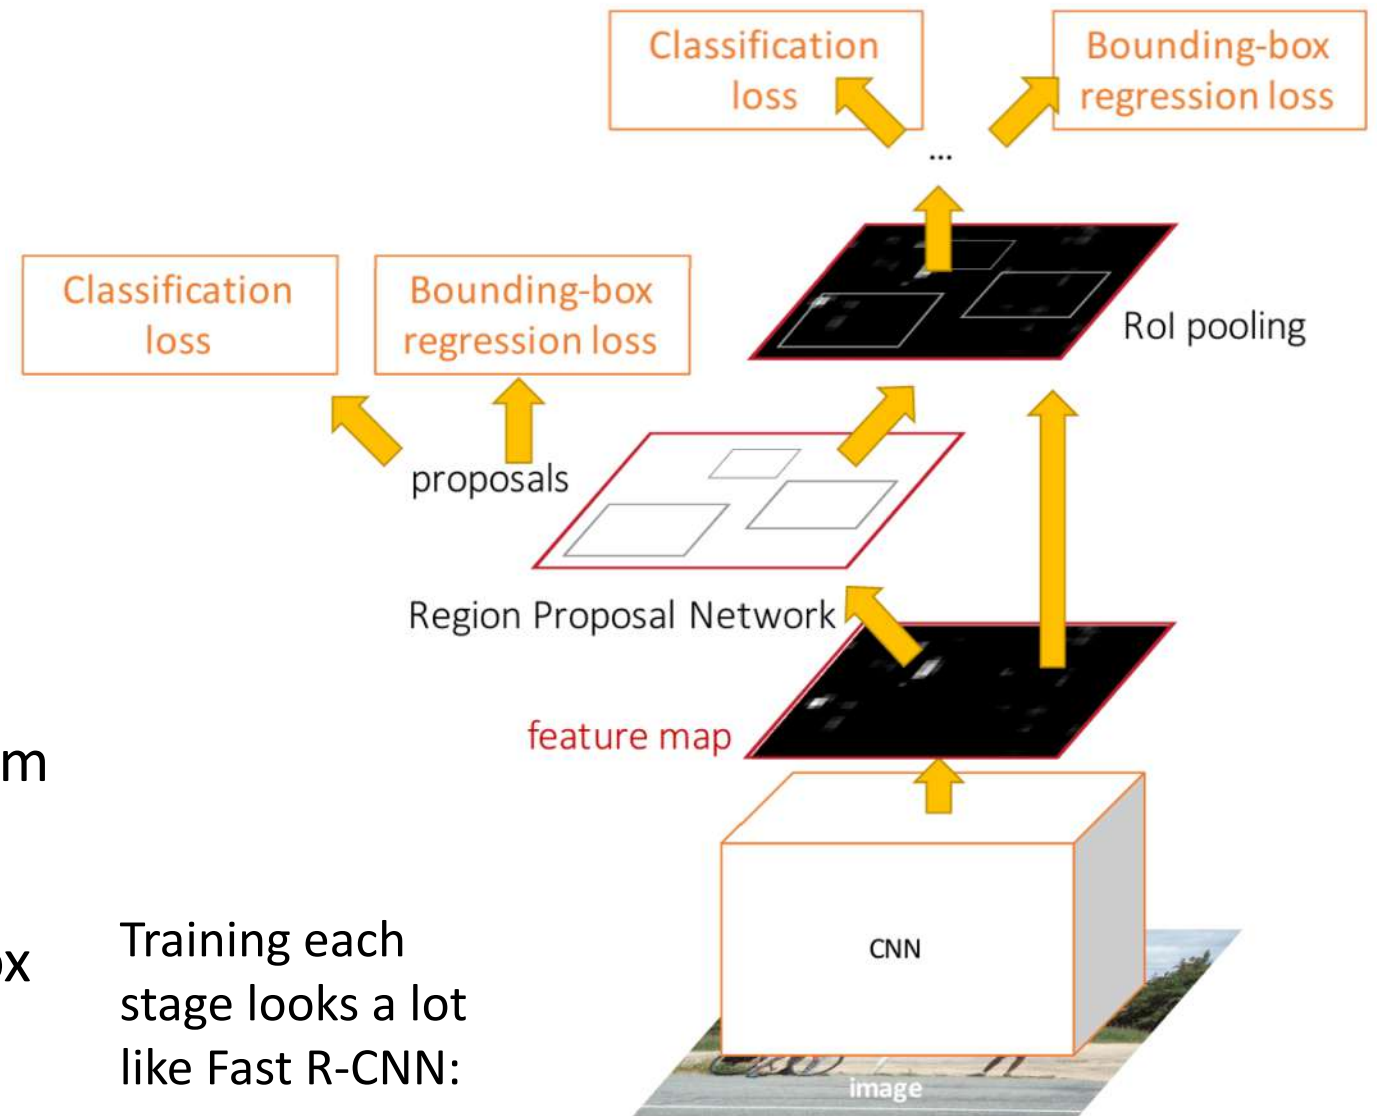
# Faster R-CNN: Learnable Region Proposals

Jointly train with 4 losses:

1. **RPN classification:** anchor box is object / not an object
2. **RPN regression:** predict transform from anchor box to proposal box
3. **Object classification:** classify proposals as background / object class
4. **Object regression:** predict transform from proposal box to object box

Anchor -> Region Proposal (Stage 1)  
Region Proposal -> Object Box (Stage 2)

Training each stage looks a lot like Fast R-CNN:






# Faster R-CNN Training: RPN Training

RPN predicts Object / Background for each anchor, as well as regresses from anchor to object box

Input Image



Backbone CNN





Image Features





RPN gives lots of **anchors** which we classify as pos / neg / neutral by matching with ground-truth

GT Boxes

Positive


Neutral

Negative




Class target: Obj

Box target: →




Class target: Obj

Box target: →



Class target: Obj

Box target: →




Class target: Background

Box target: None

[This image is CC0 public domain](#)

# Faster R-CNN Training: Stage 2

Input Image



Backbone  
CNN


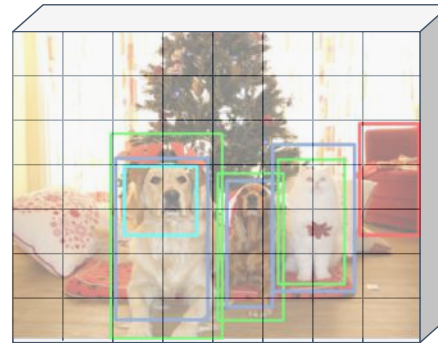


Image Features



Now proposals come from RPN rather than selective search, but otherwise this works the same as Fast R-CNN training



GT Boxes

Positive

Neutral



Negative

Crop features for each proposal, use them to predict class and box targets per region




Class target: Dog

Box target: →





Class target: Cat

Box target: →



Class target: Dog

Box target: →



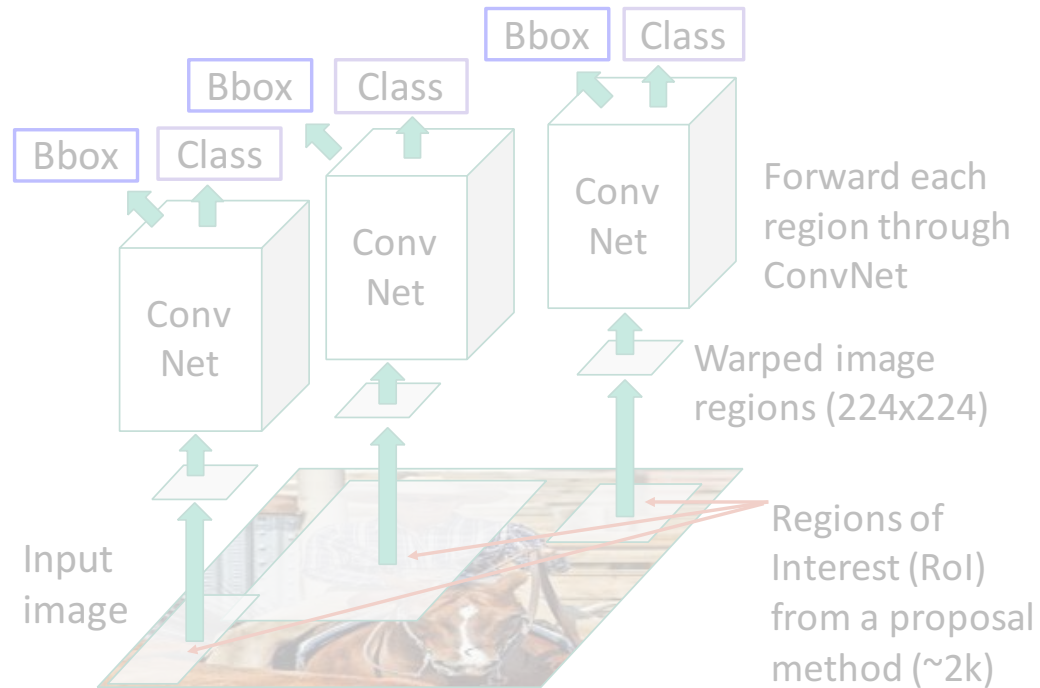
Class target: Background

Box target: None

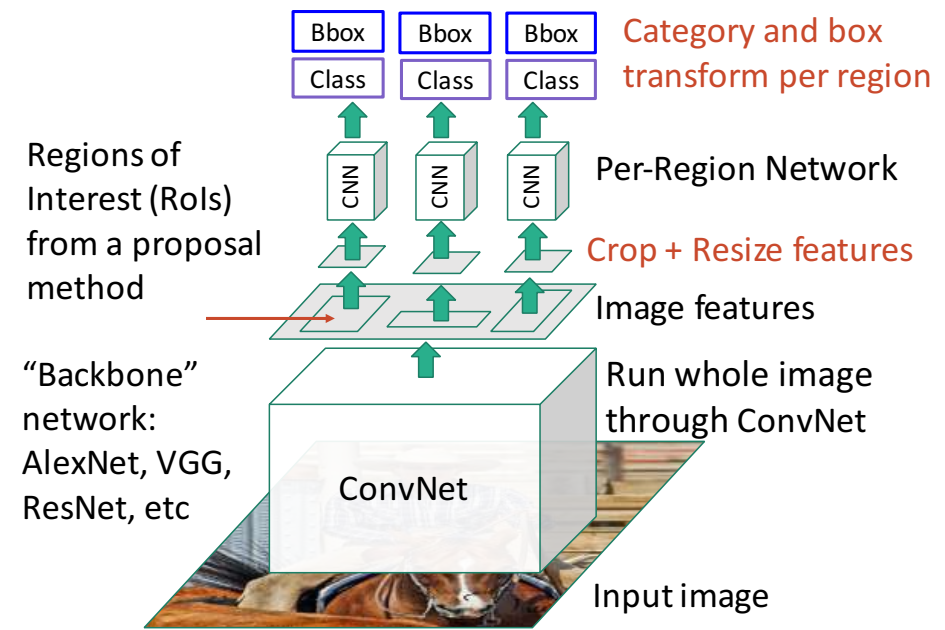
[This image is CC0 public domain](#)

# Recap: Fast R-CNN Feature Cropping

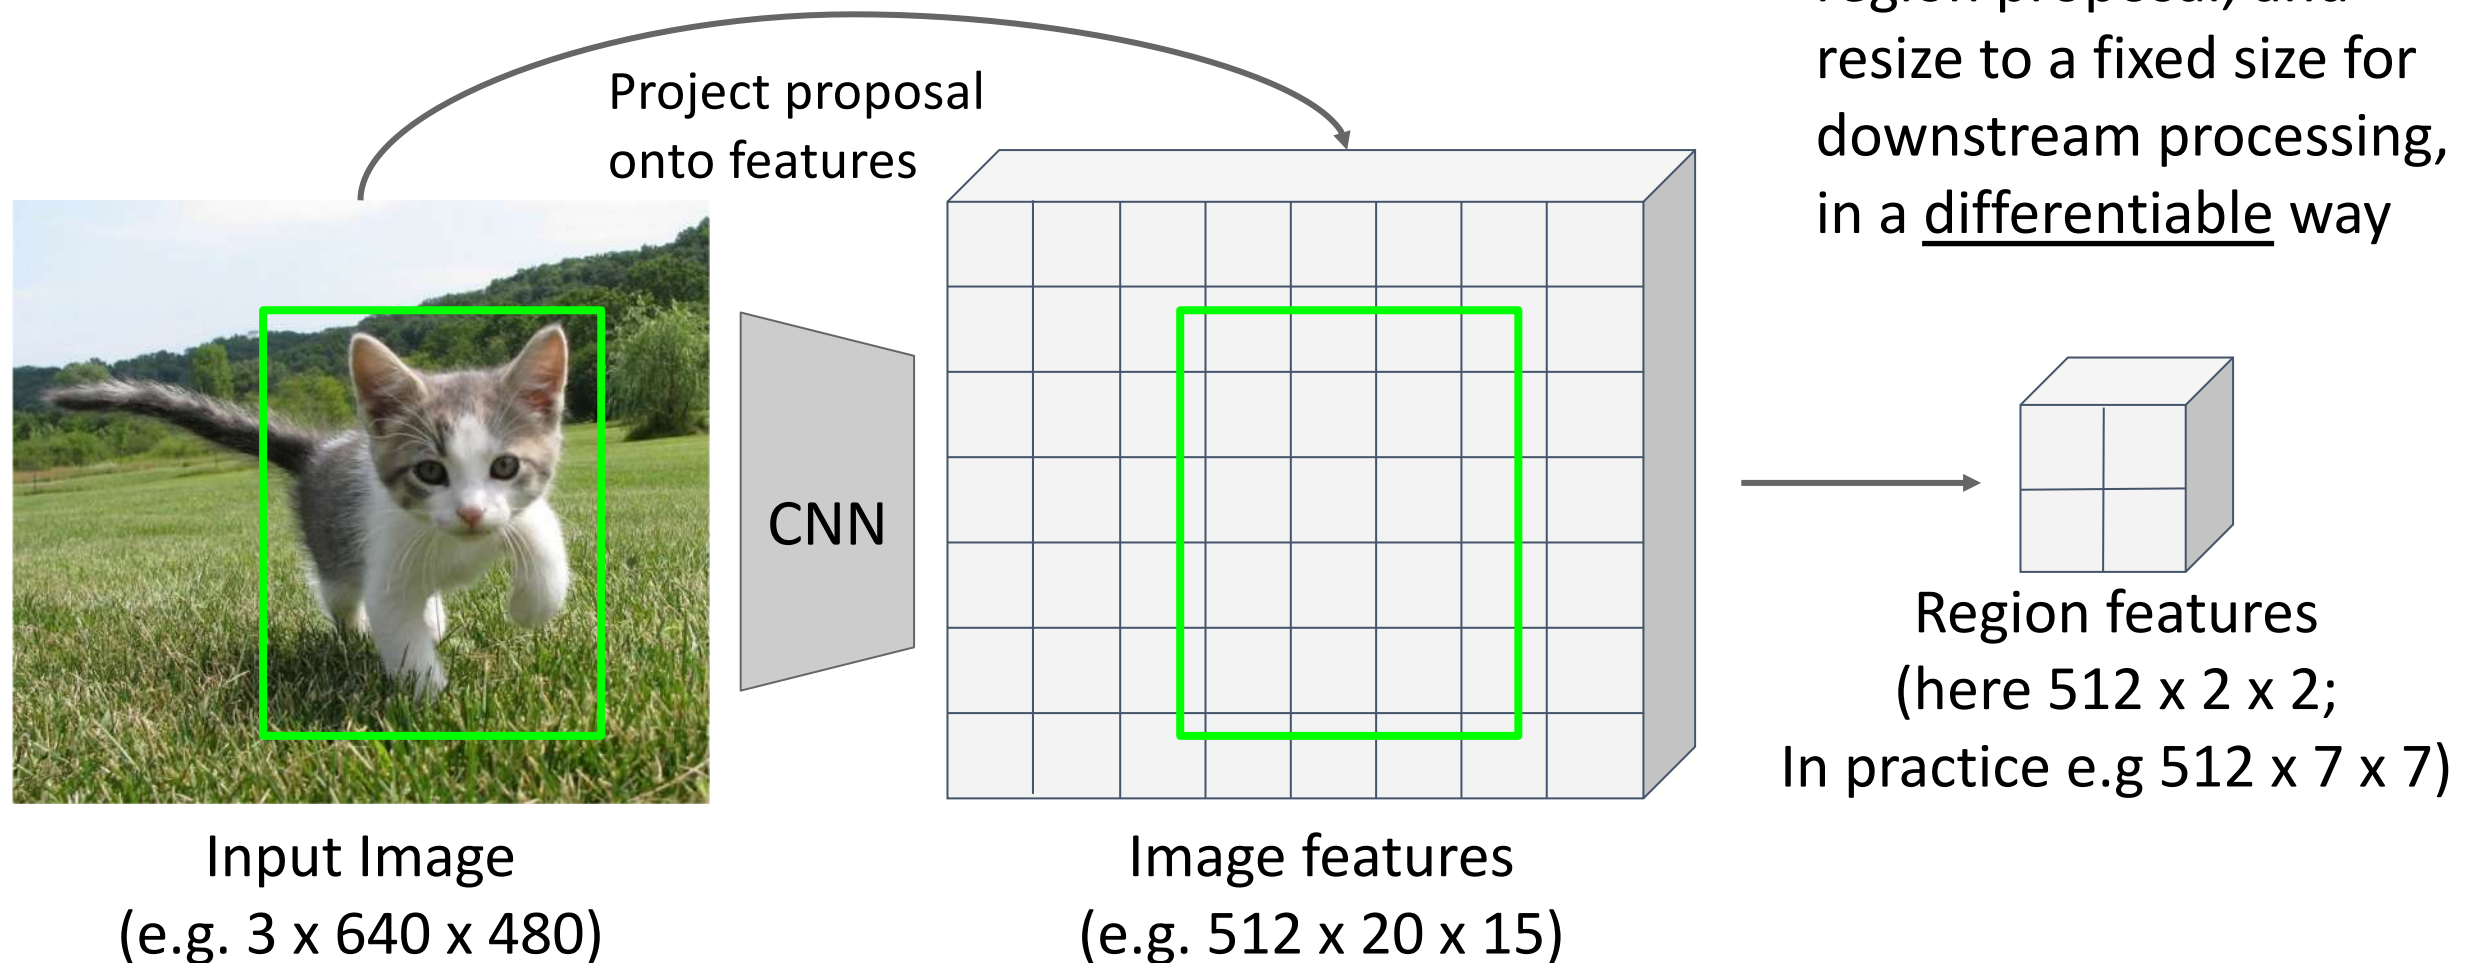
**“Slow” R-CNN:** Run CNN independently for each region



**Fast R-CNN:** Apply differentiable cropping to shared image features



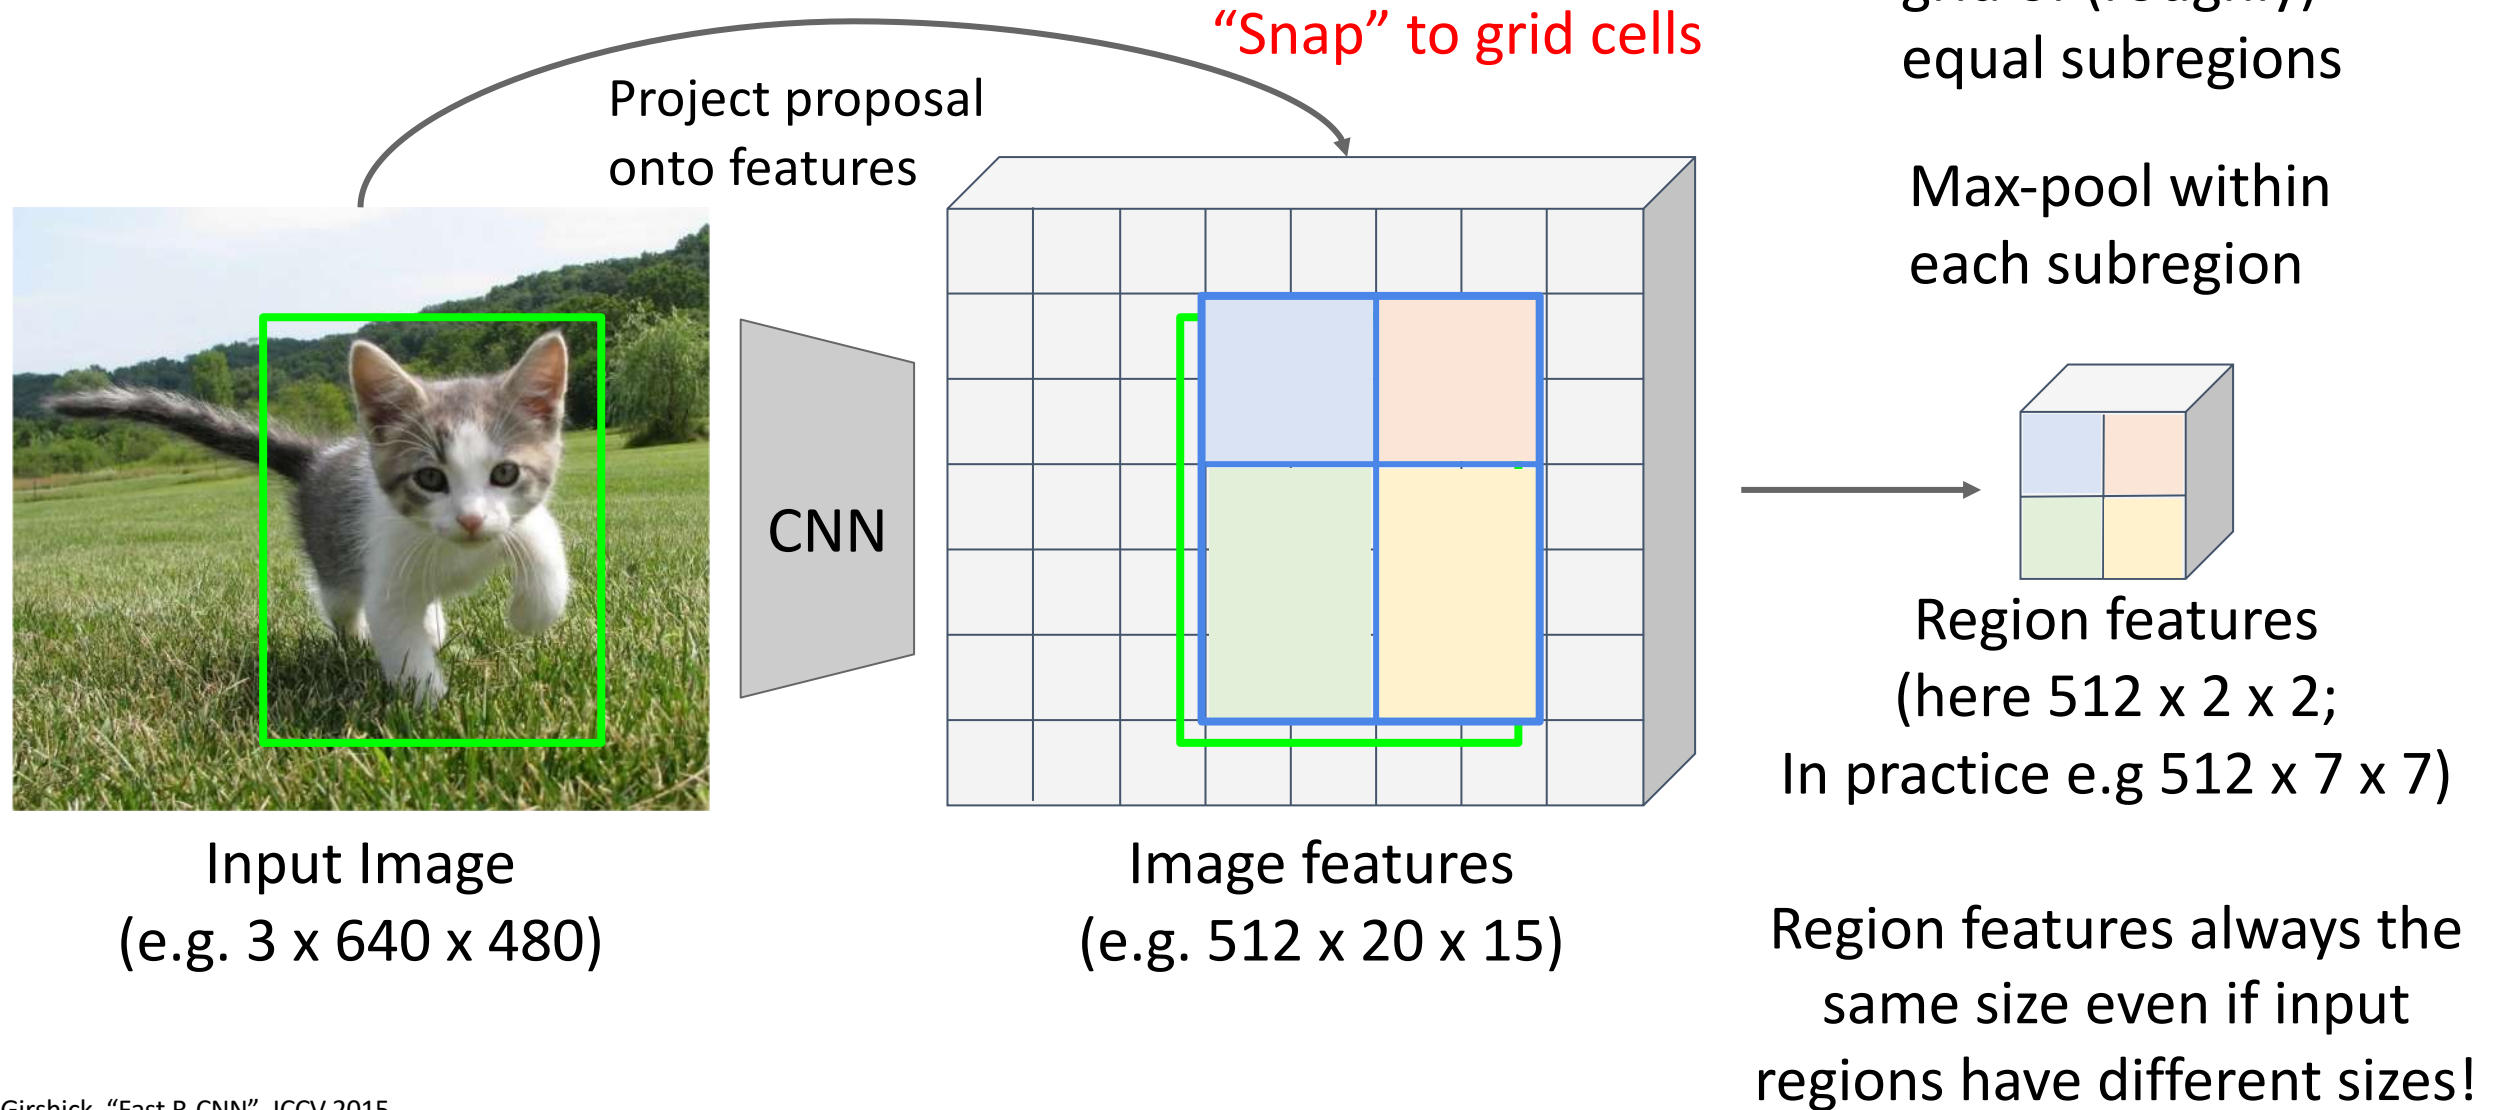
# Cropping Features



Girshick, "Fast R-CNN", ICCV 2015.

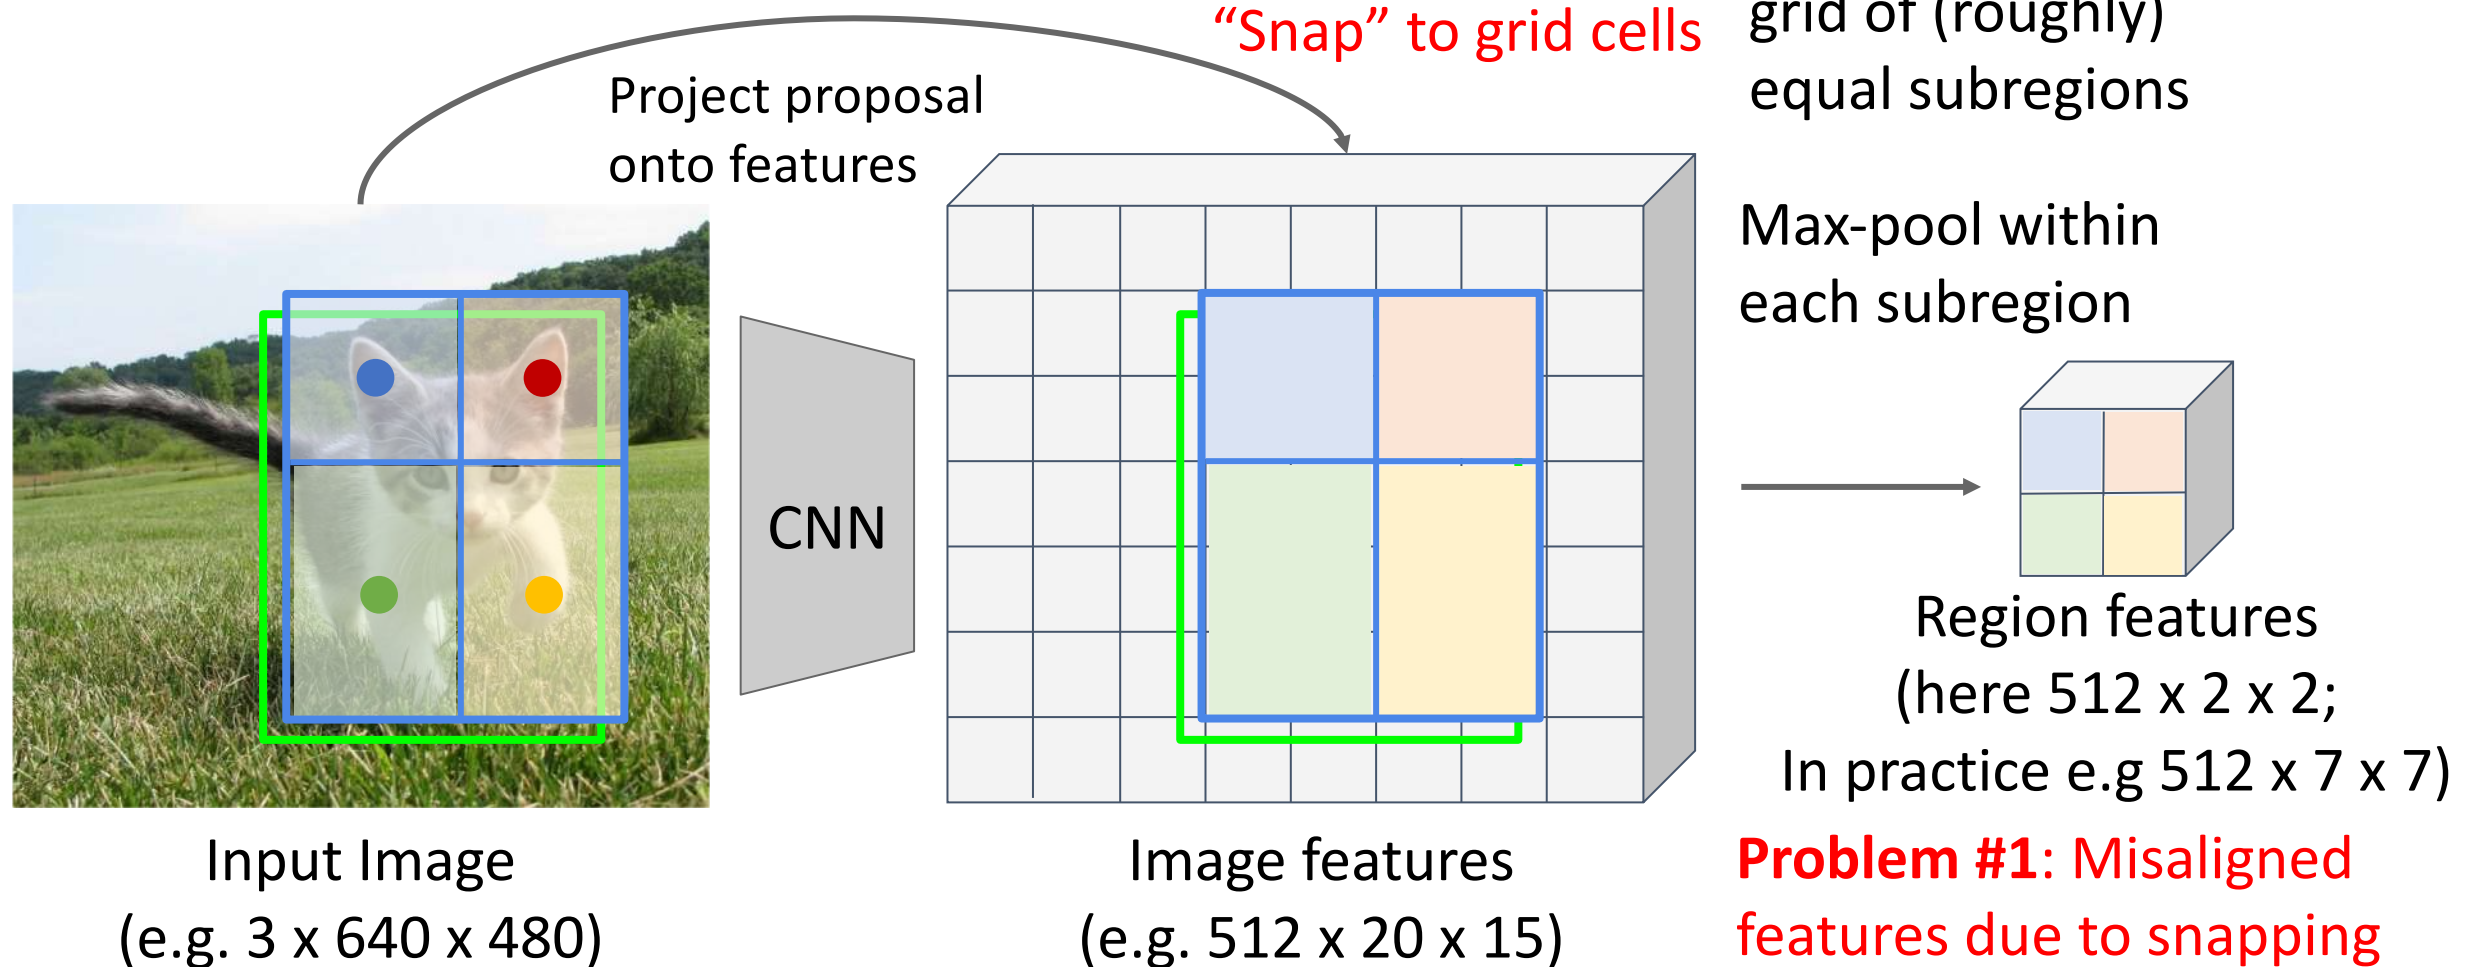


# Cropping Features: RoI Pool



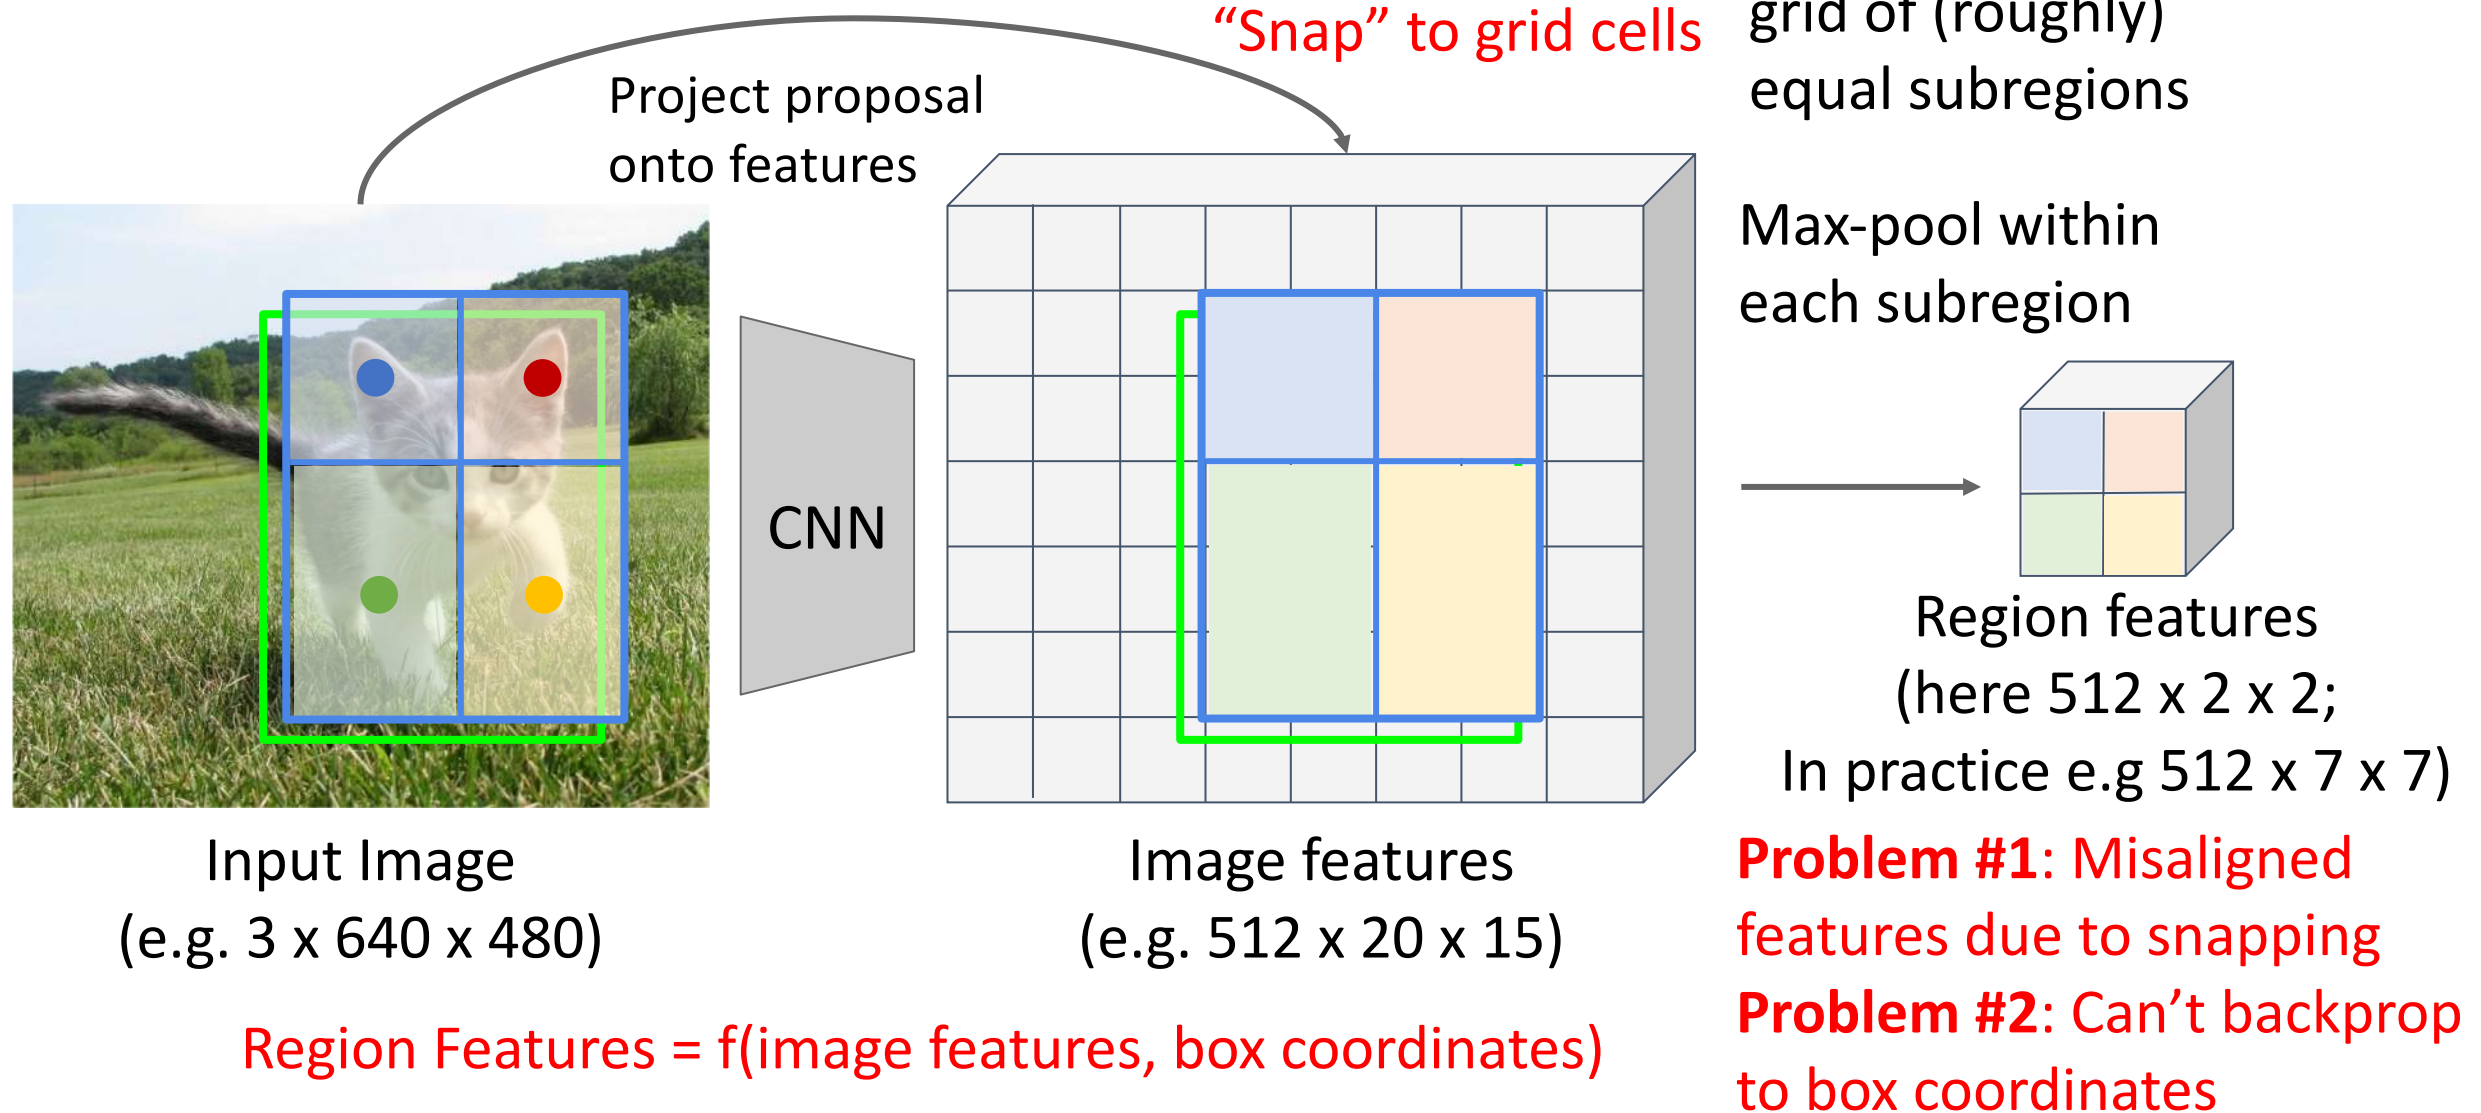
Girshick, “Fast R-CNN”, ICCV 2015.

# Cropping Features: RoI Pool



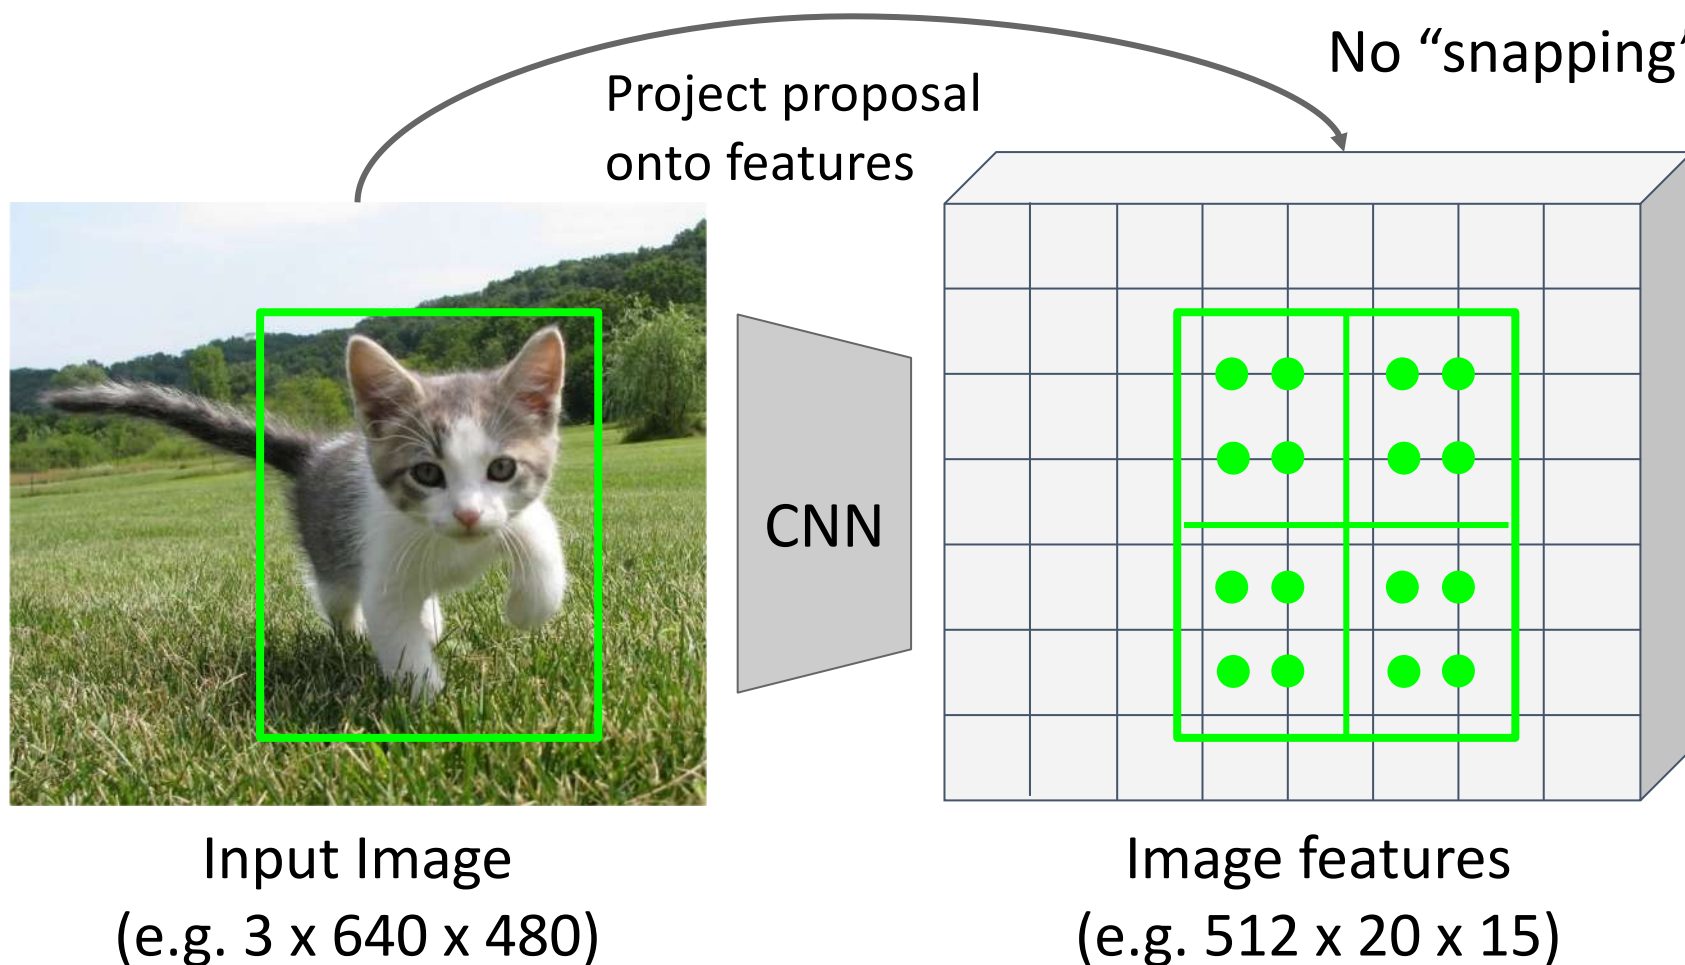


# Cropping Features: RoI Pool



# Cropping Features: RoI Align

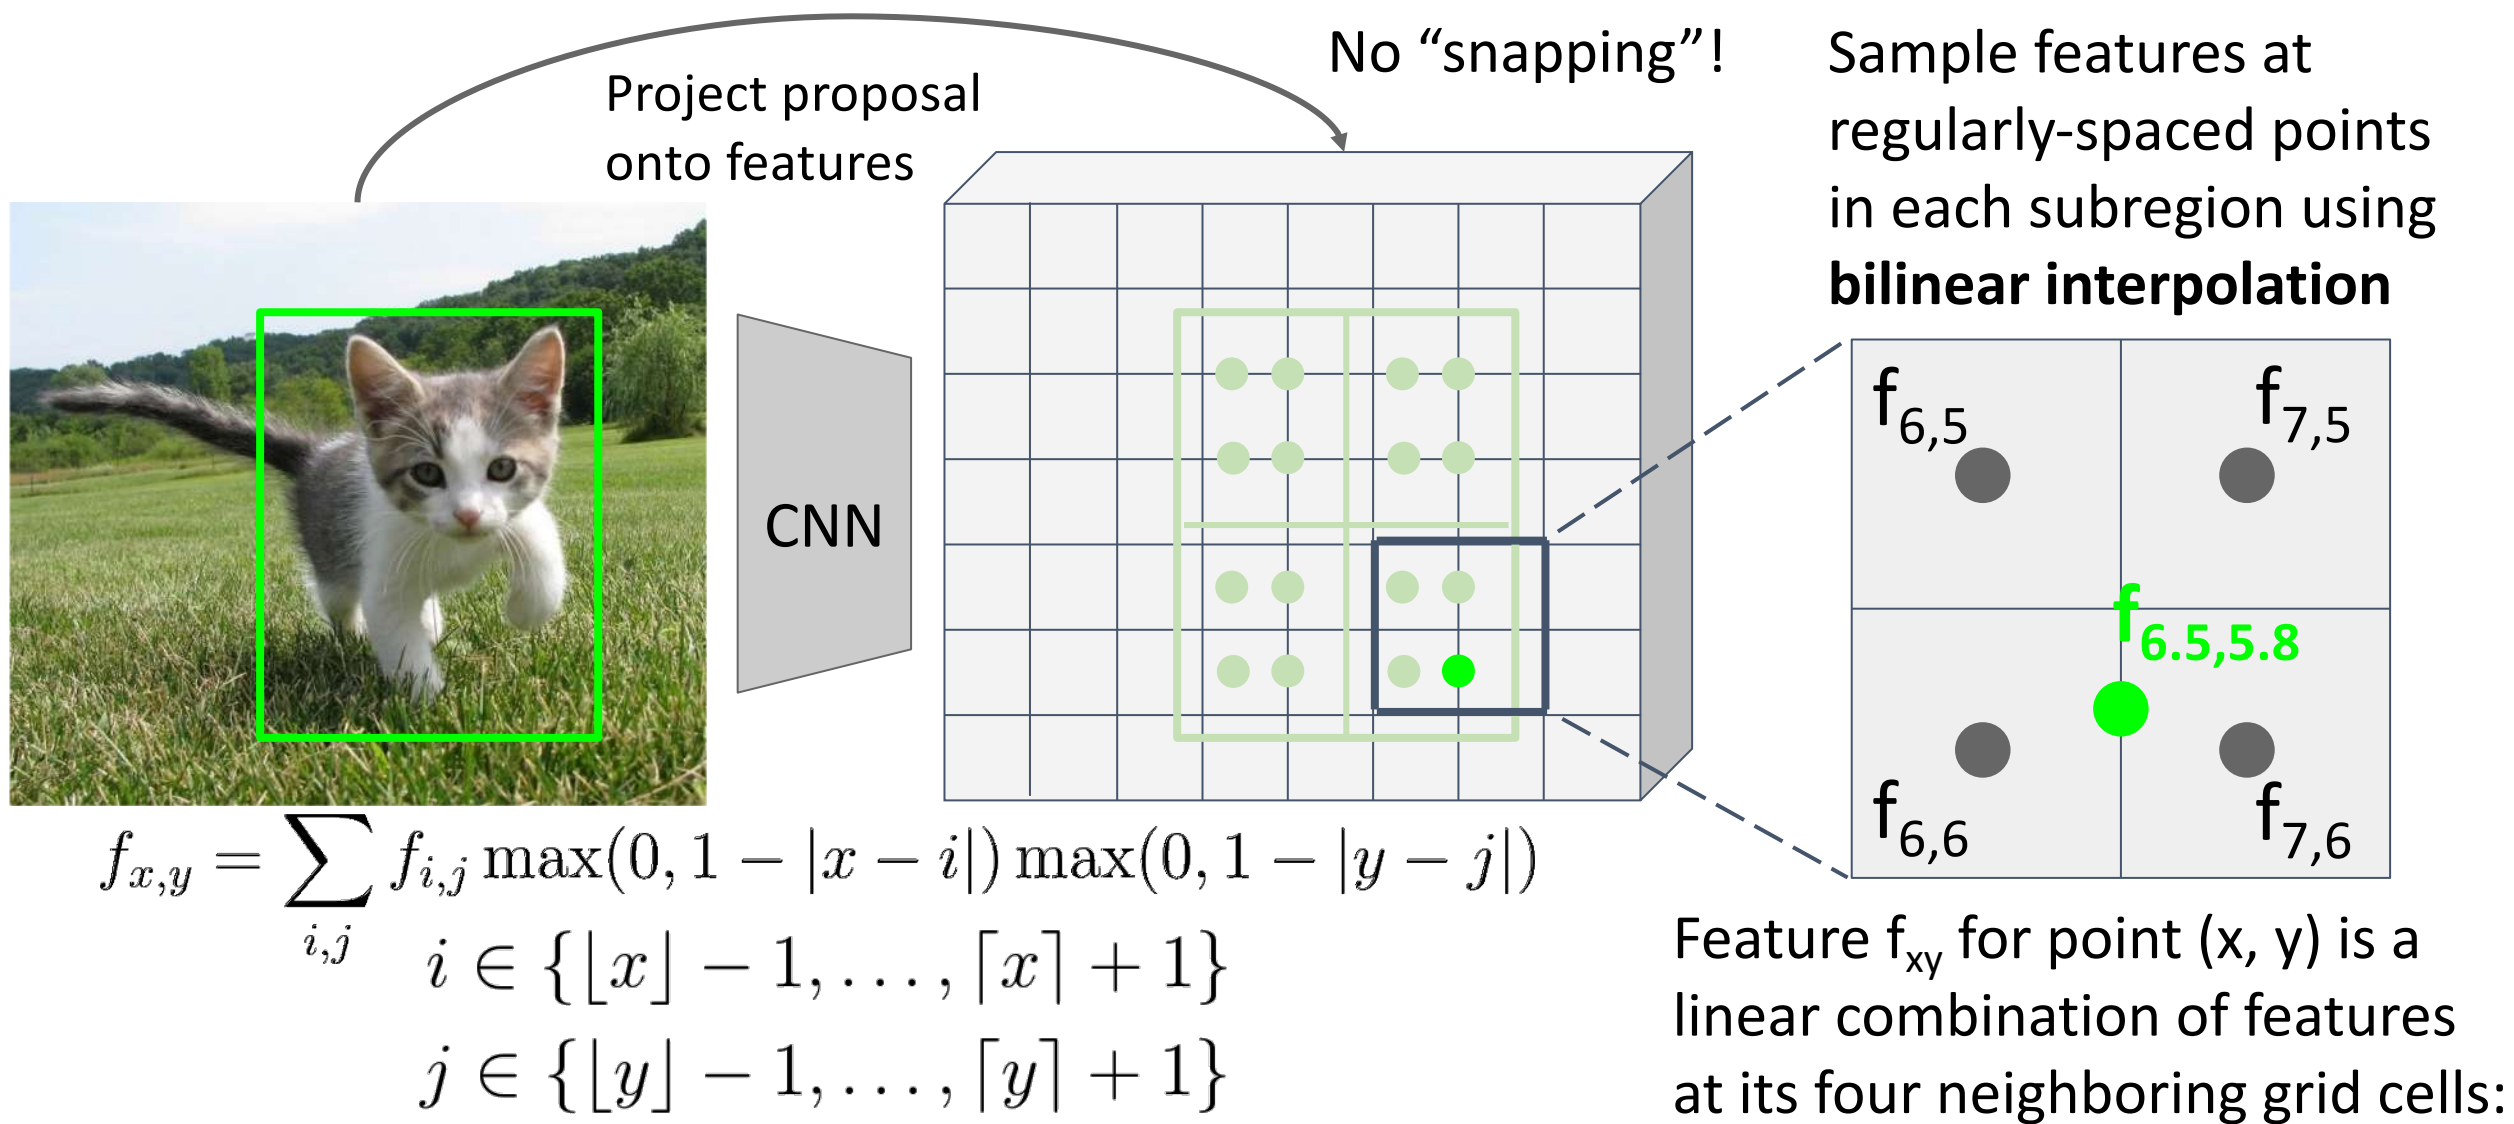
Divide into equal-sized subregions  
(may not be aligned to grid!)



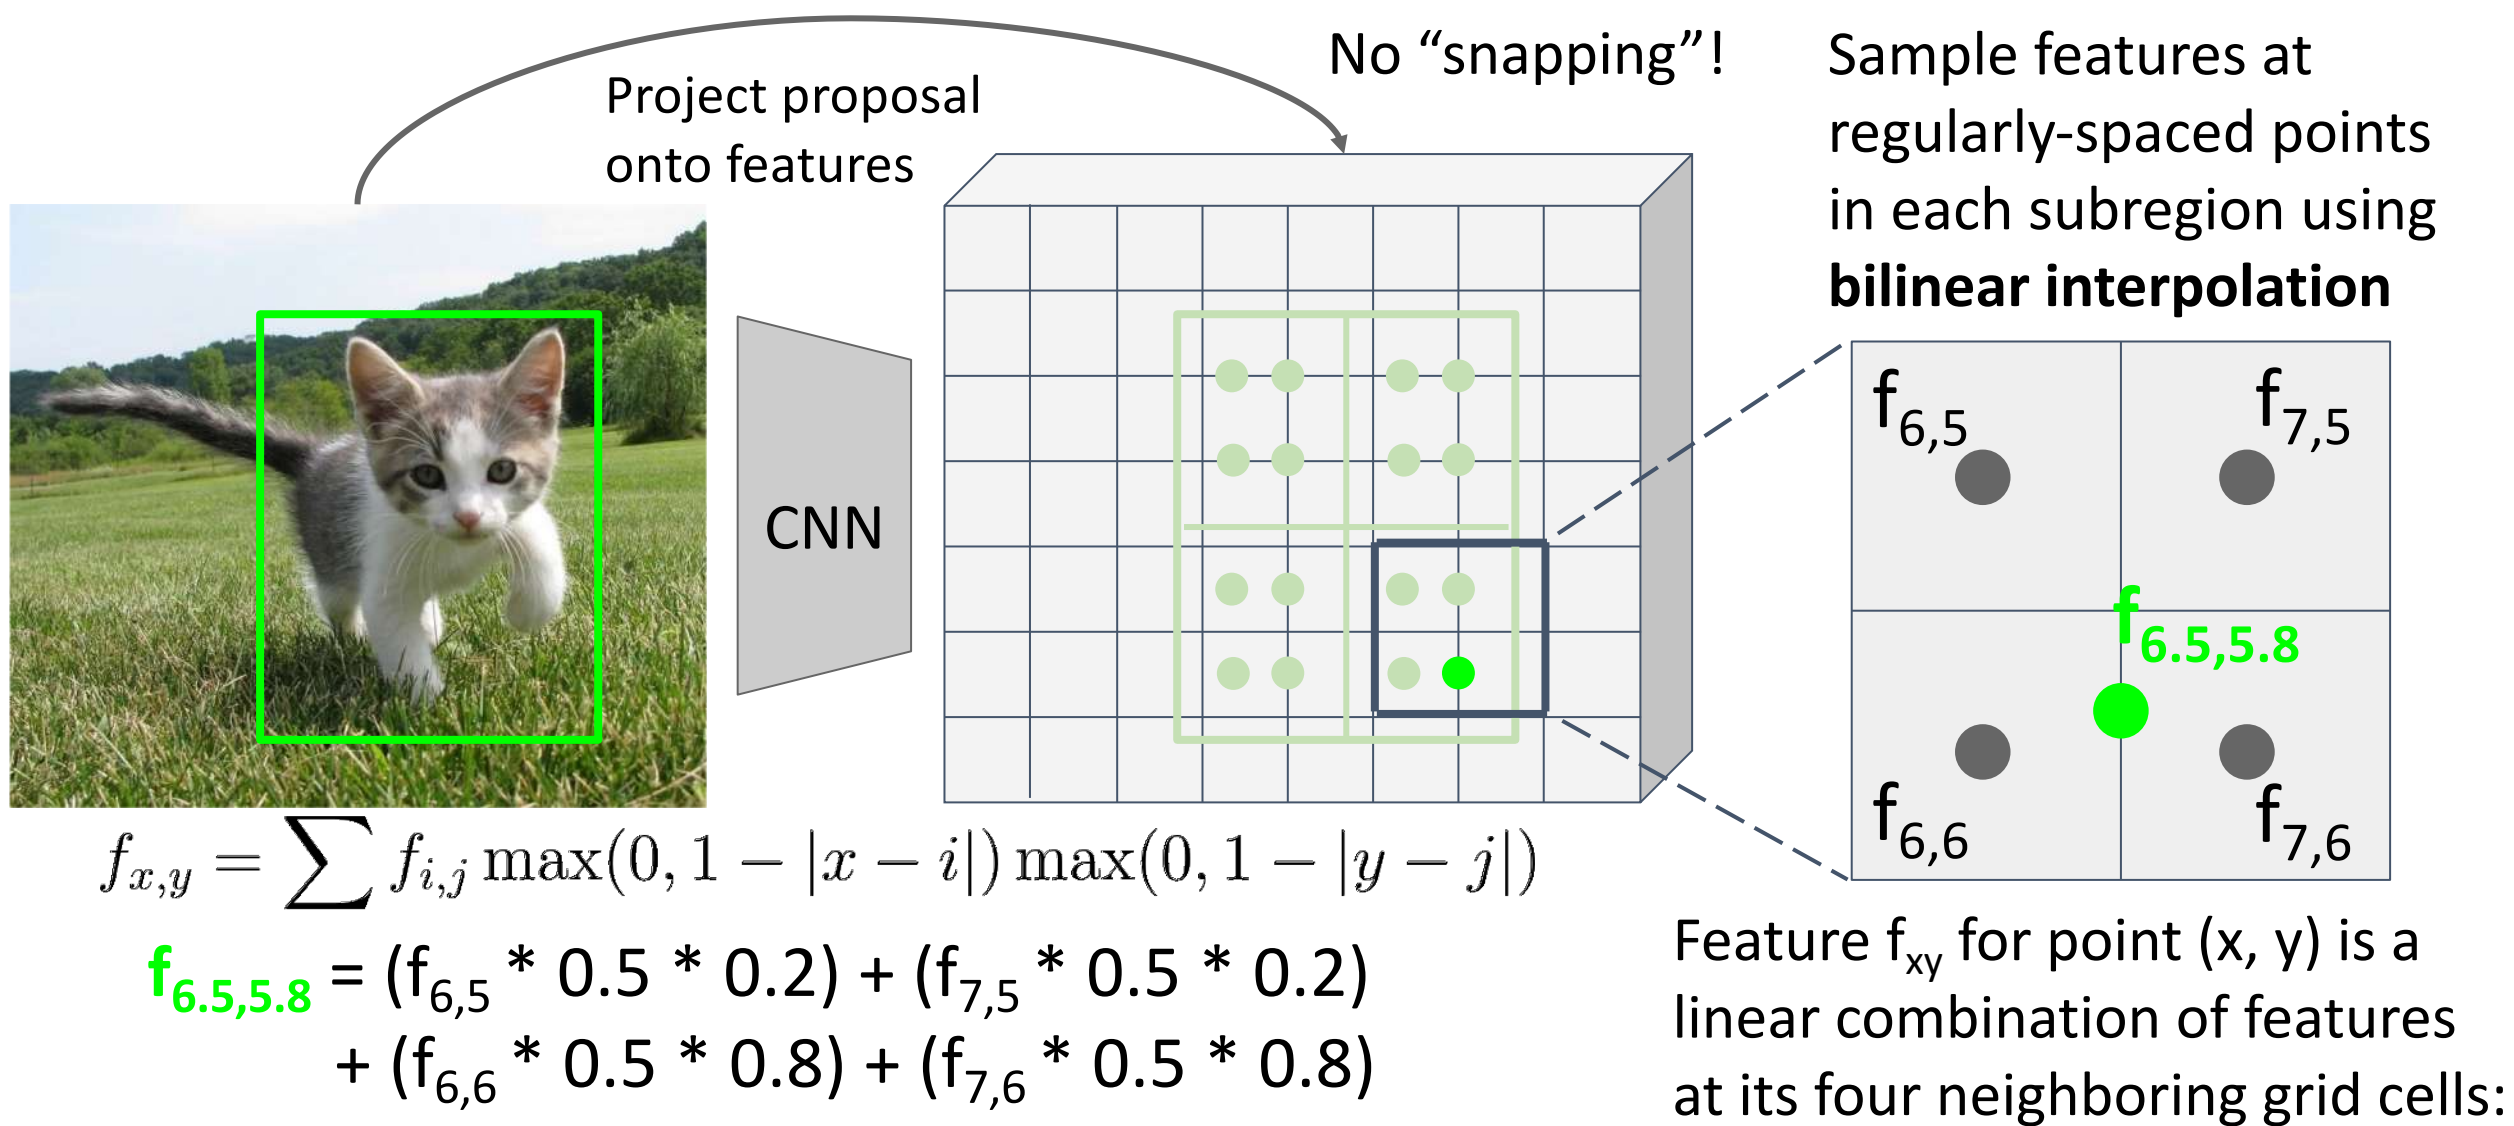
Sample features at  
regularly-spaced points  
in each subregion using  
**bilinear interpolation**

# Cropping Features: RoI Align

Divide into equal-sized subregions  
(may not be aligned to grid!)



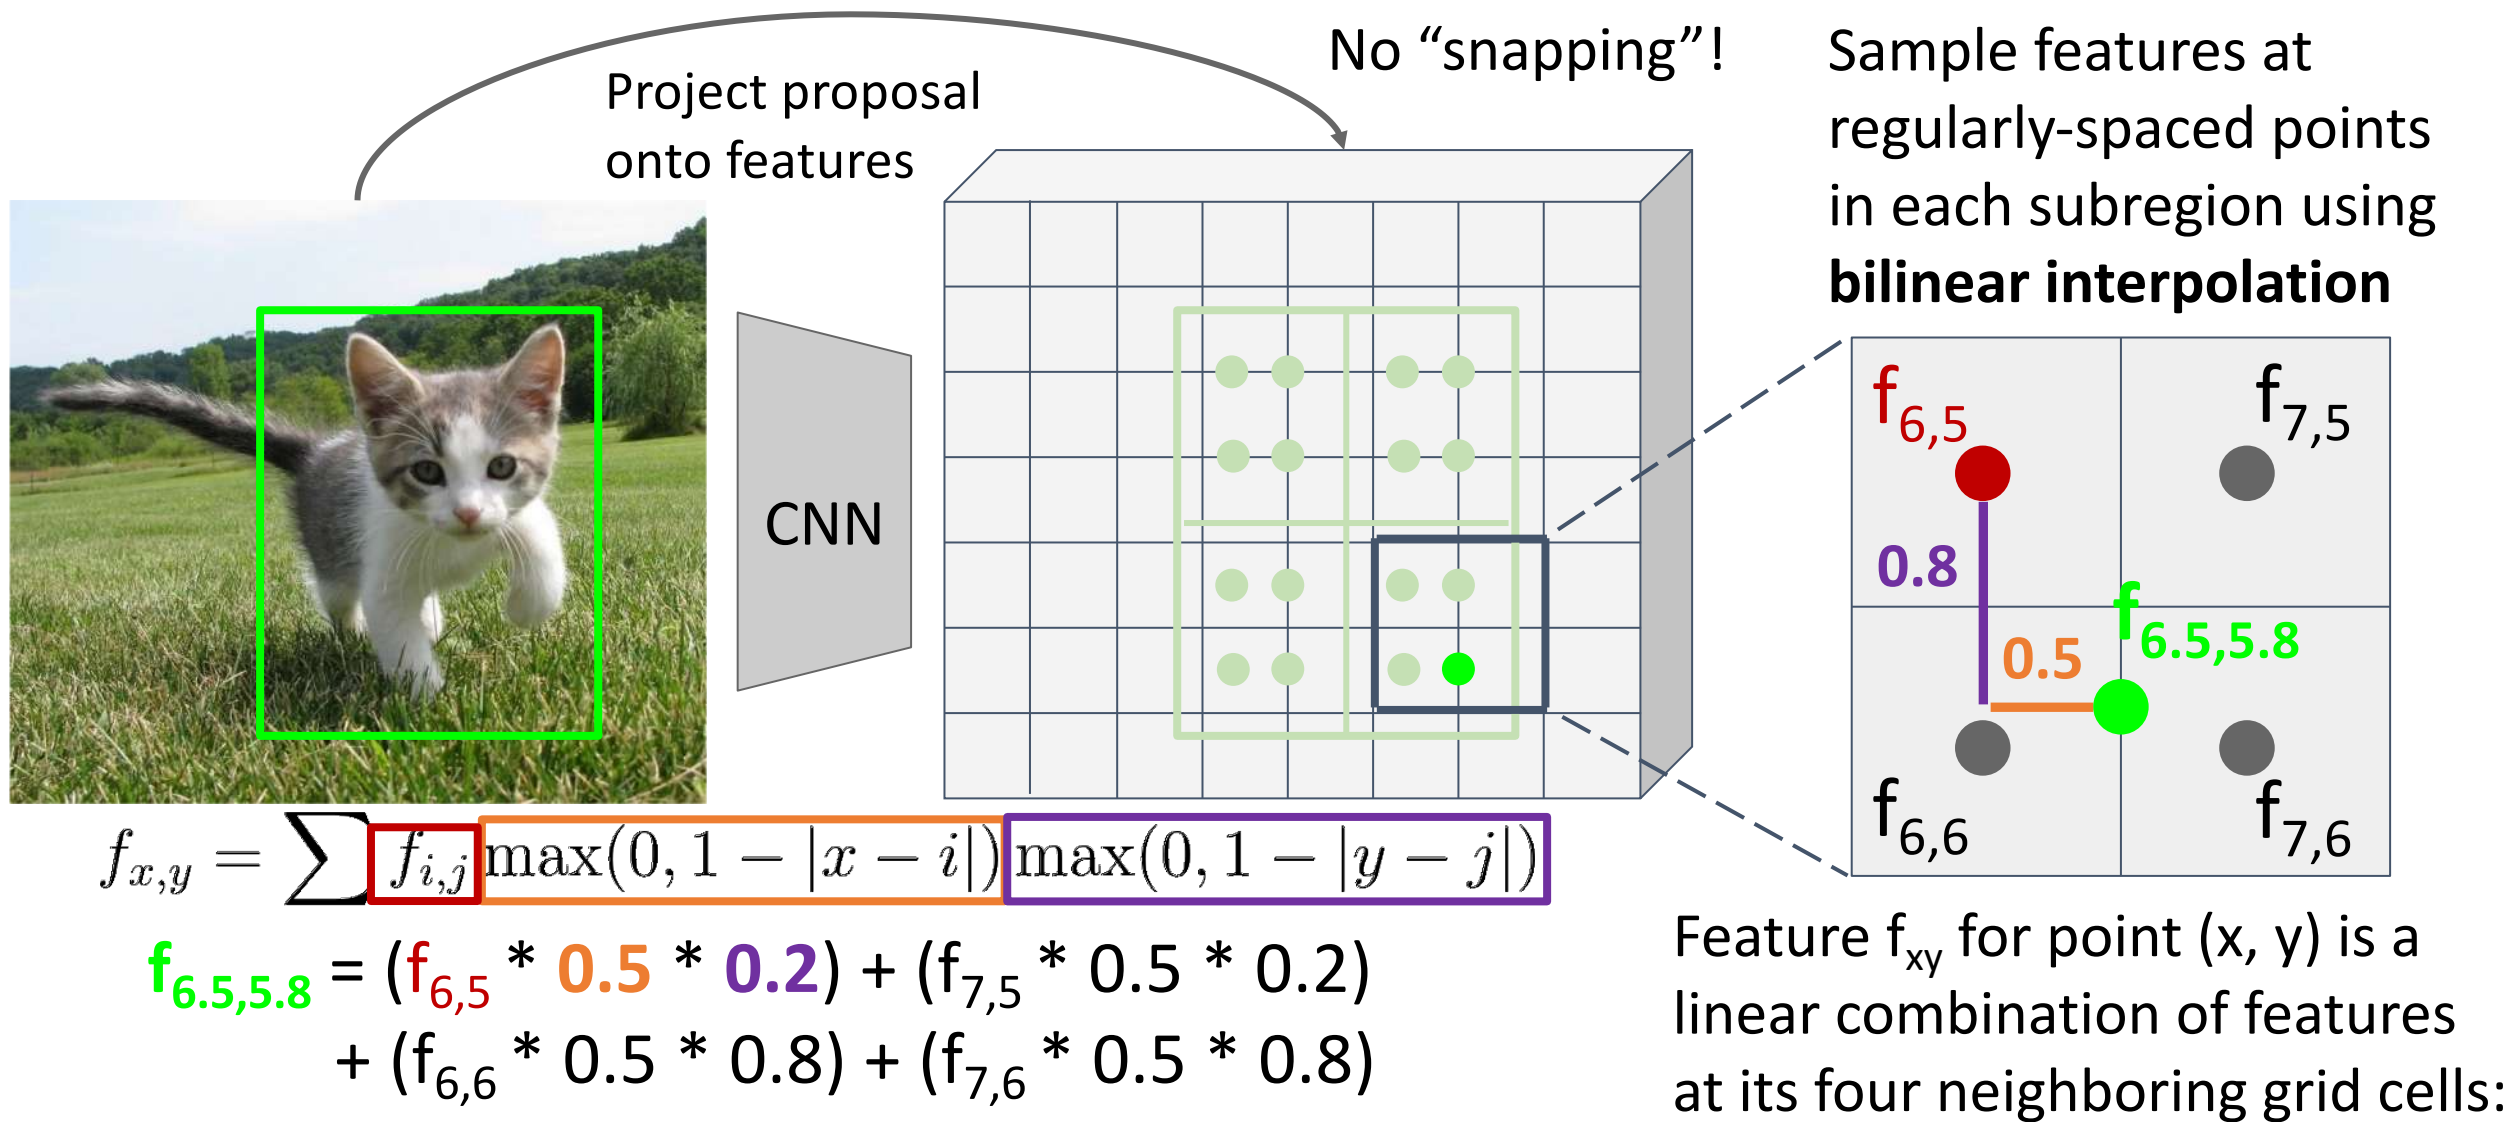
# Cropping Features: RoI Align



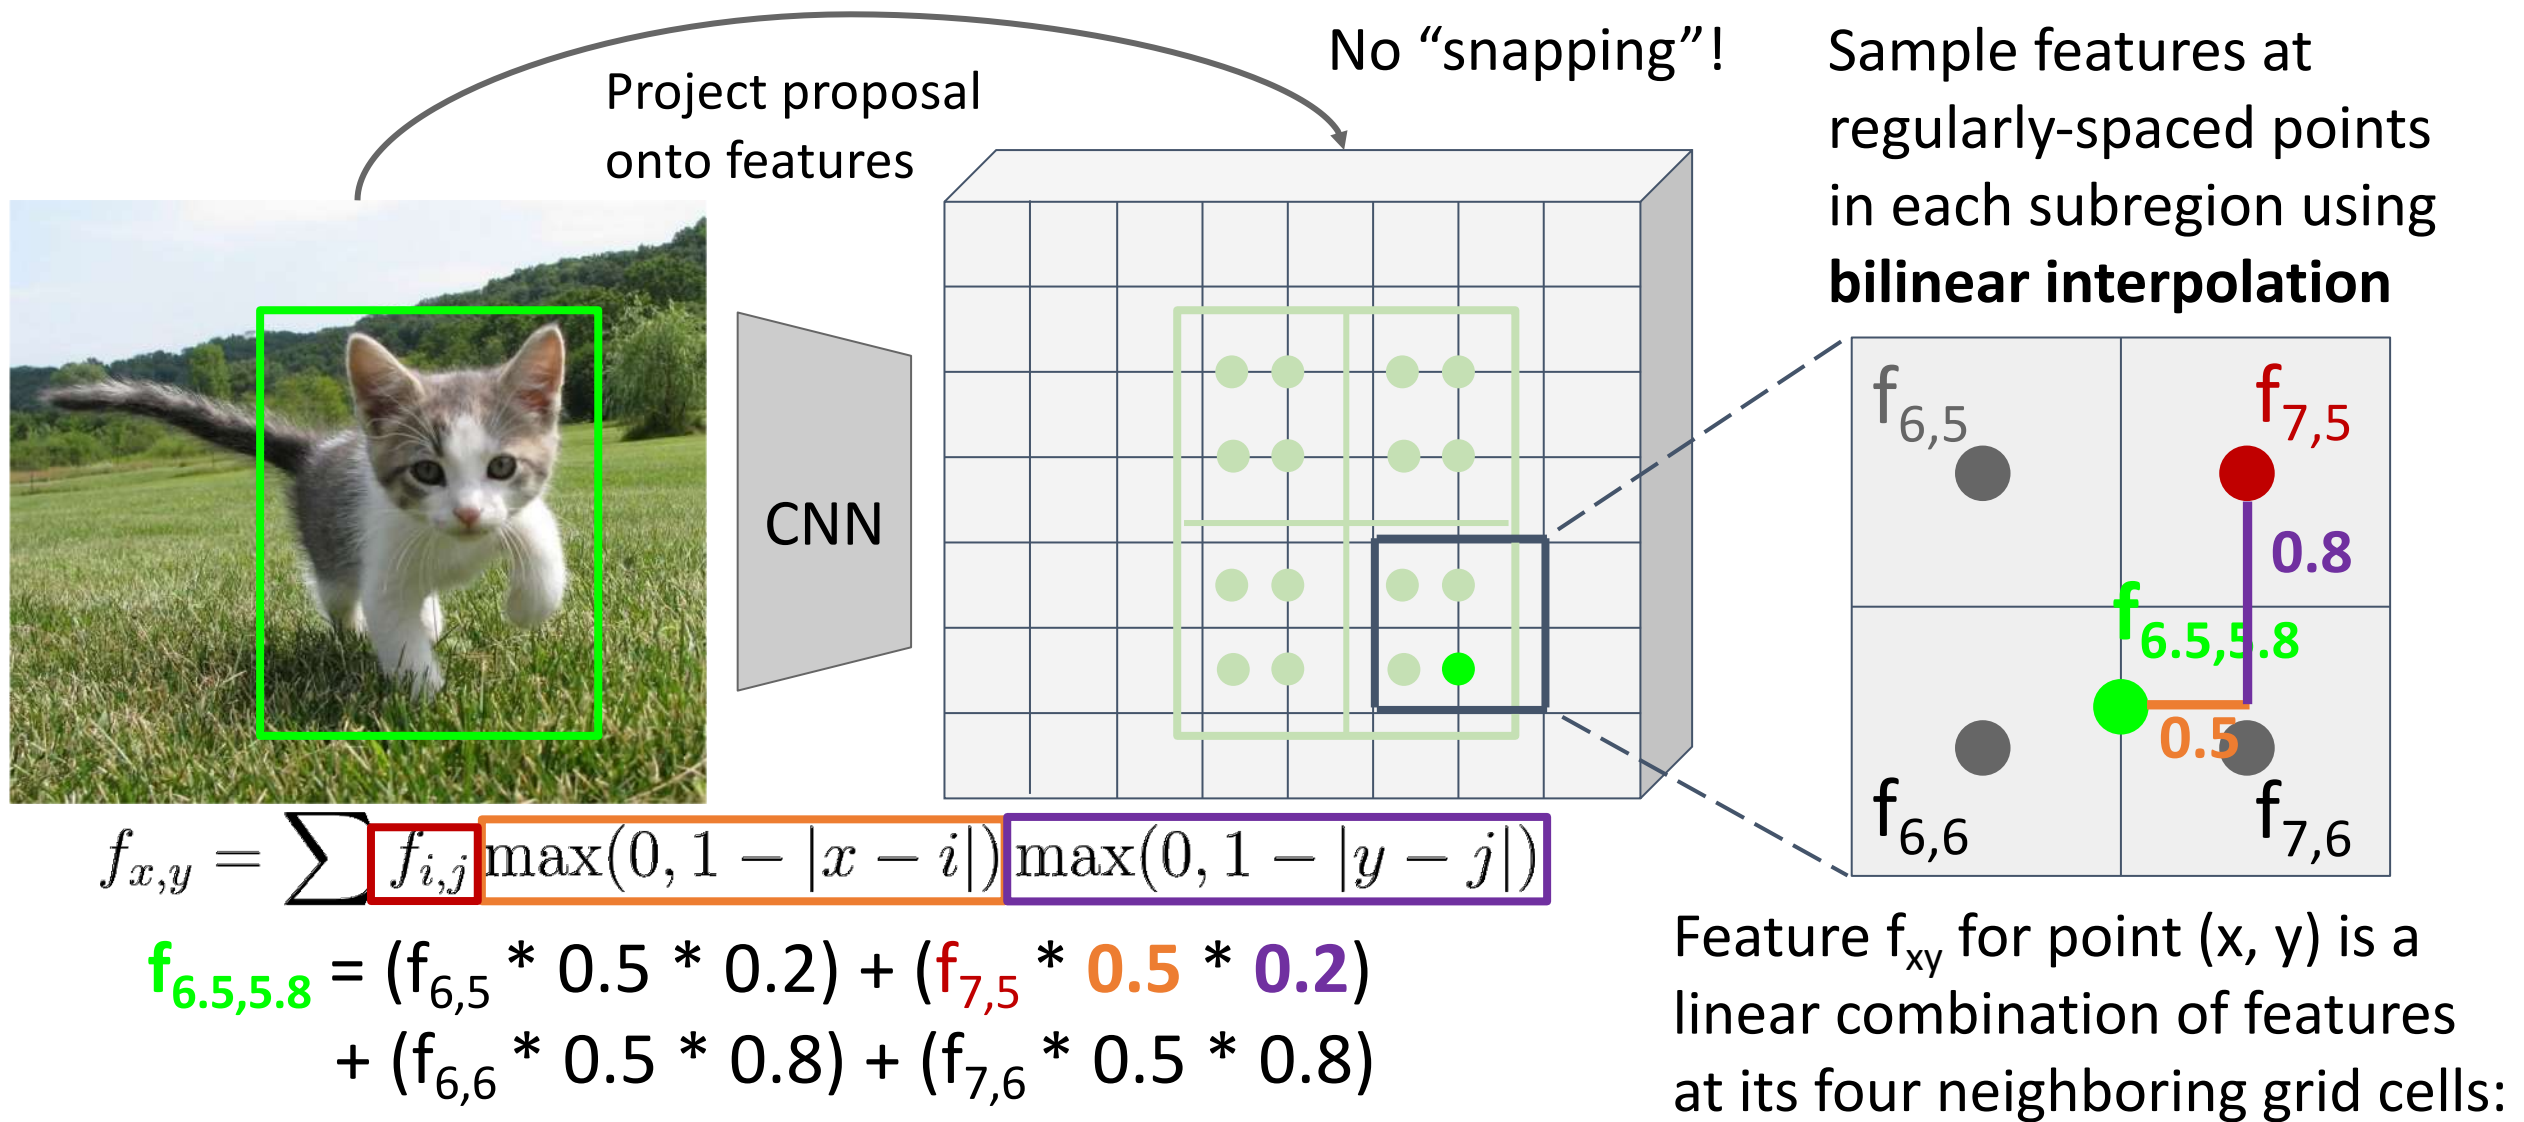


# Cropping Features: RoI Align

Divide into equal-sized subregions  
(may not be aligned to grid!)



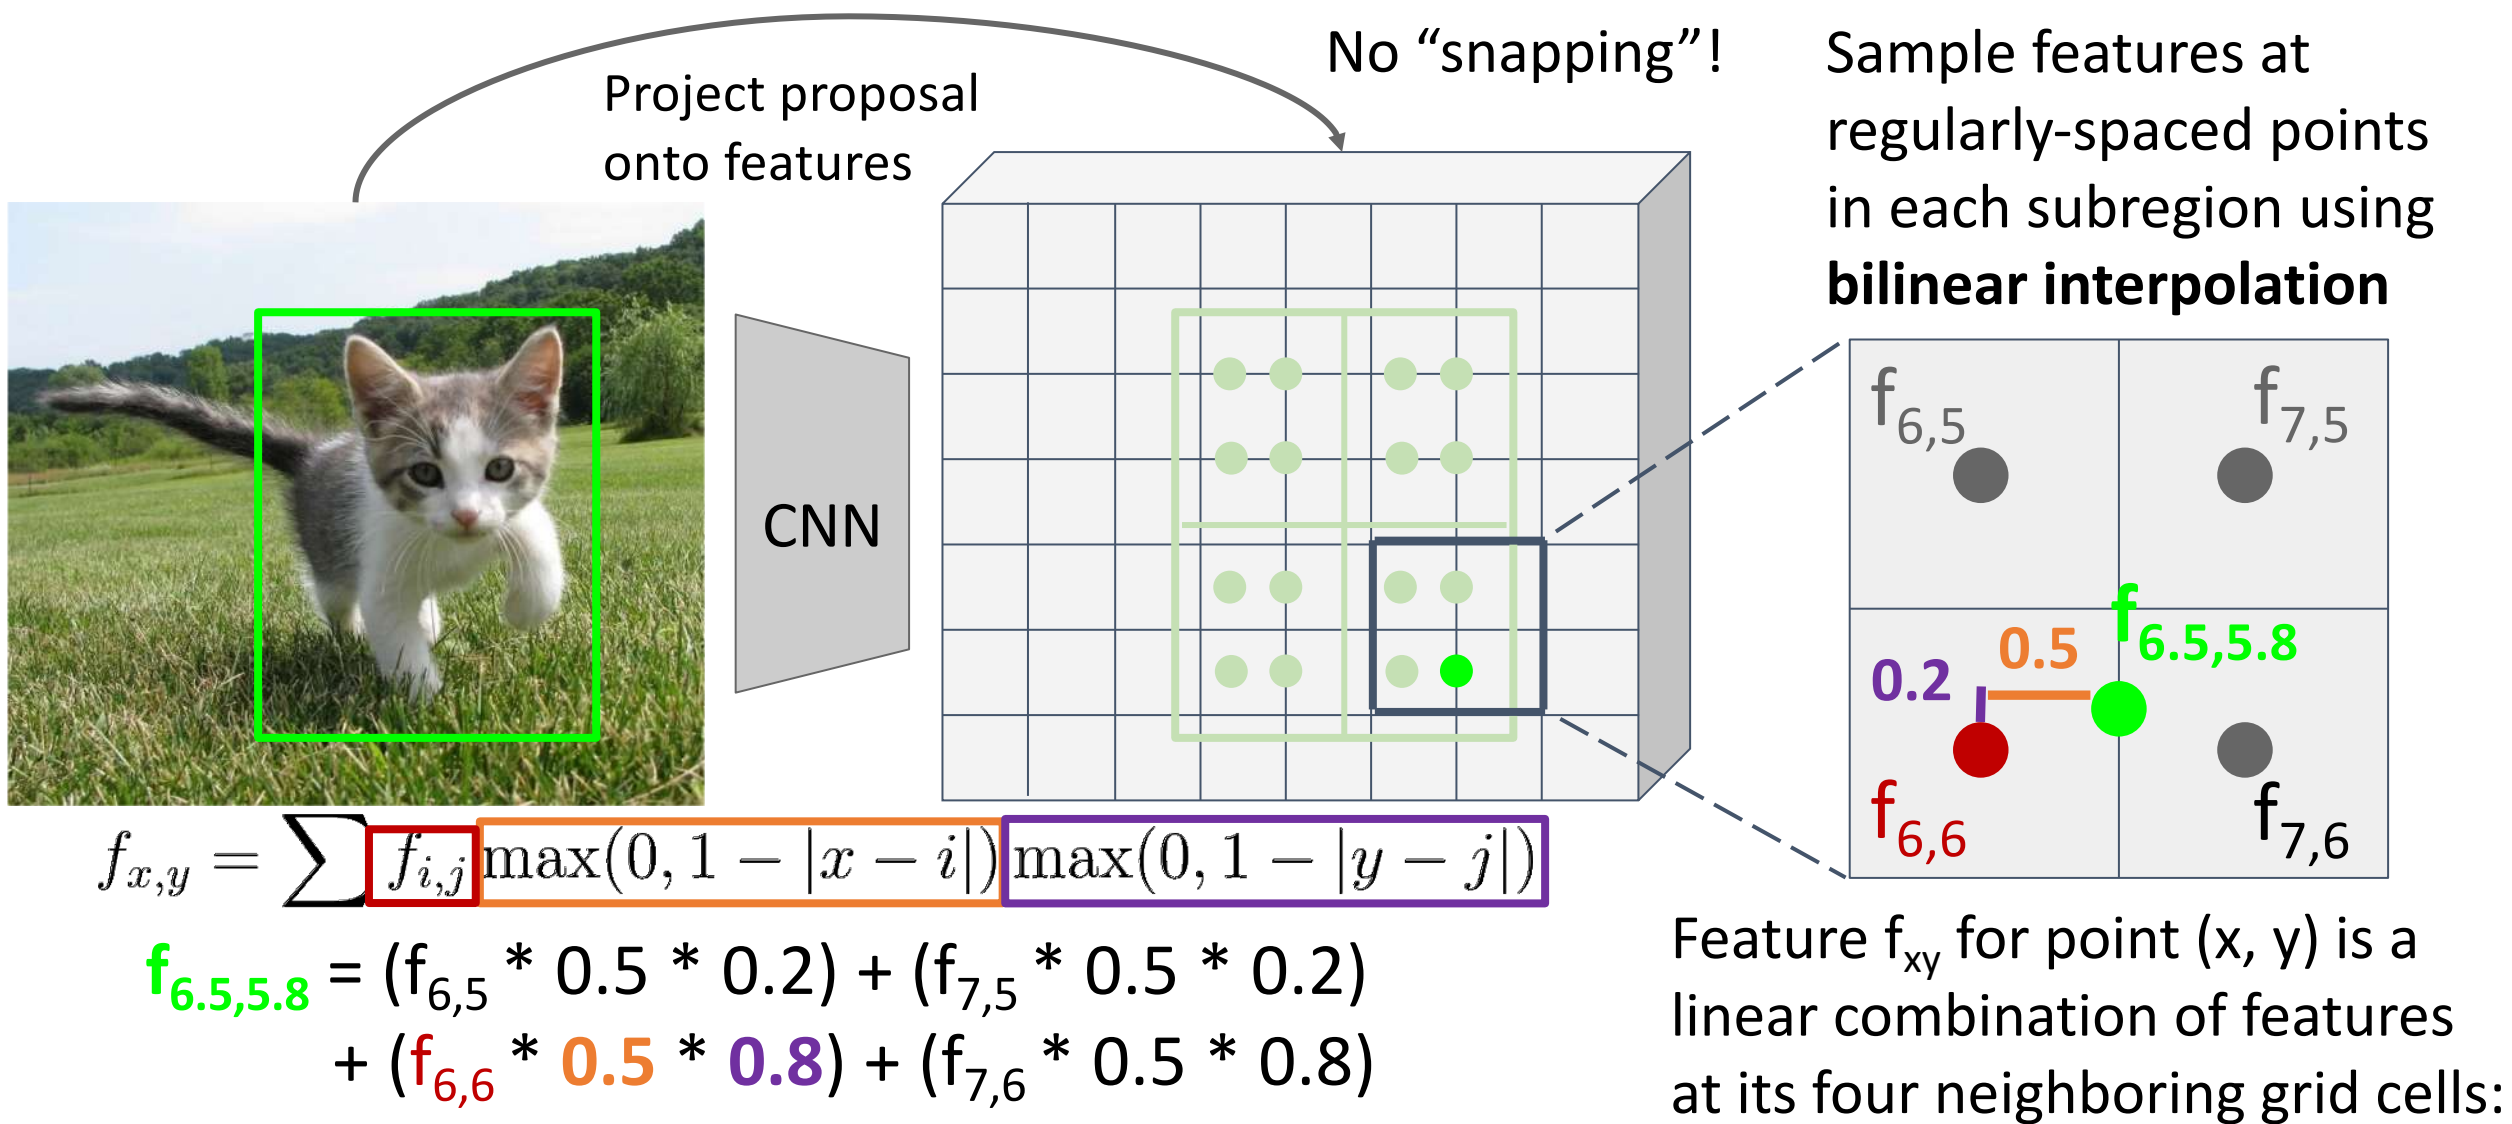
# Cropping Features: RoI Align



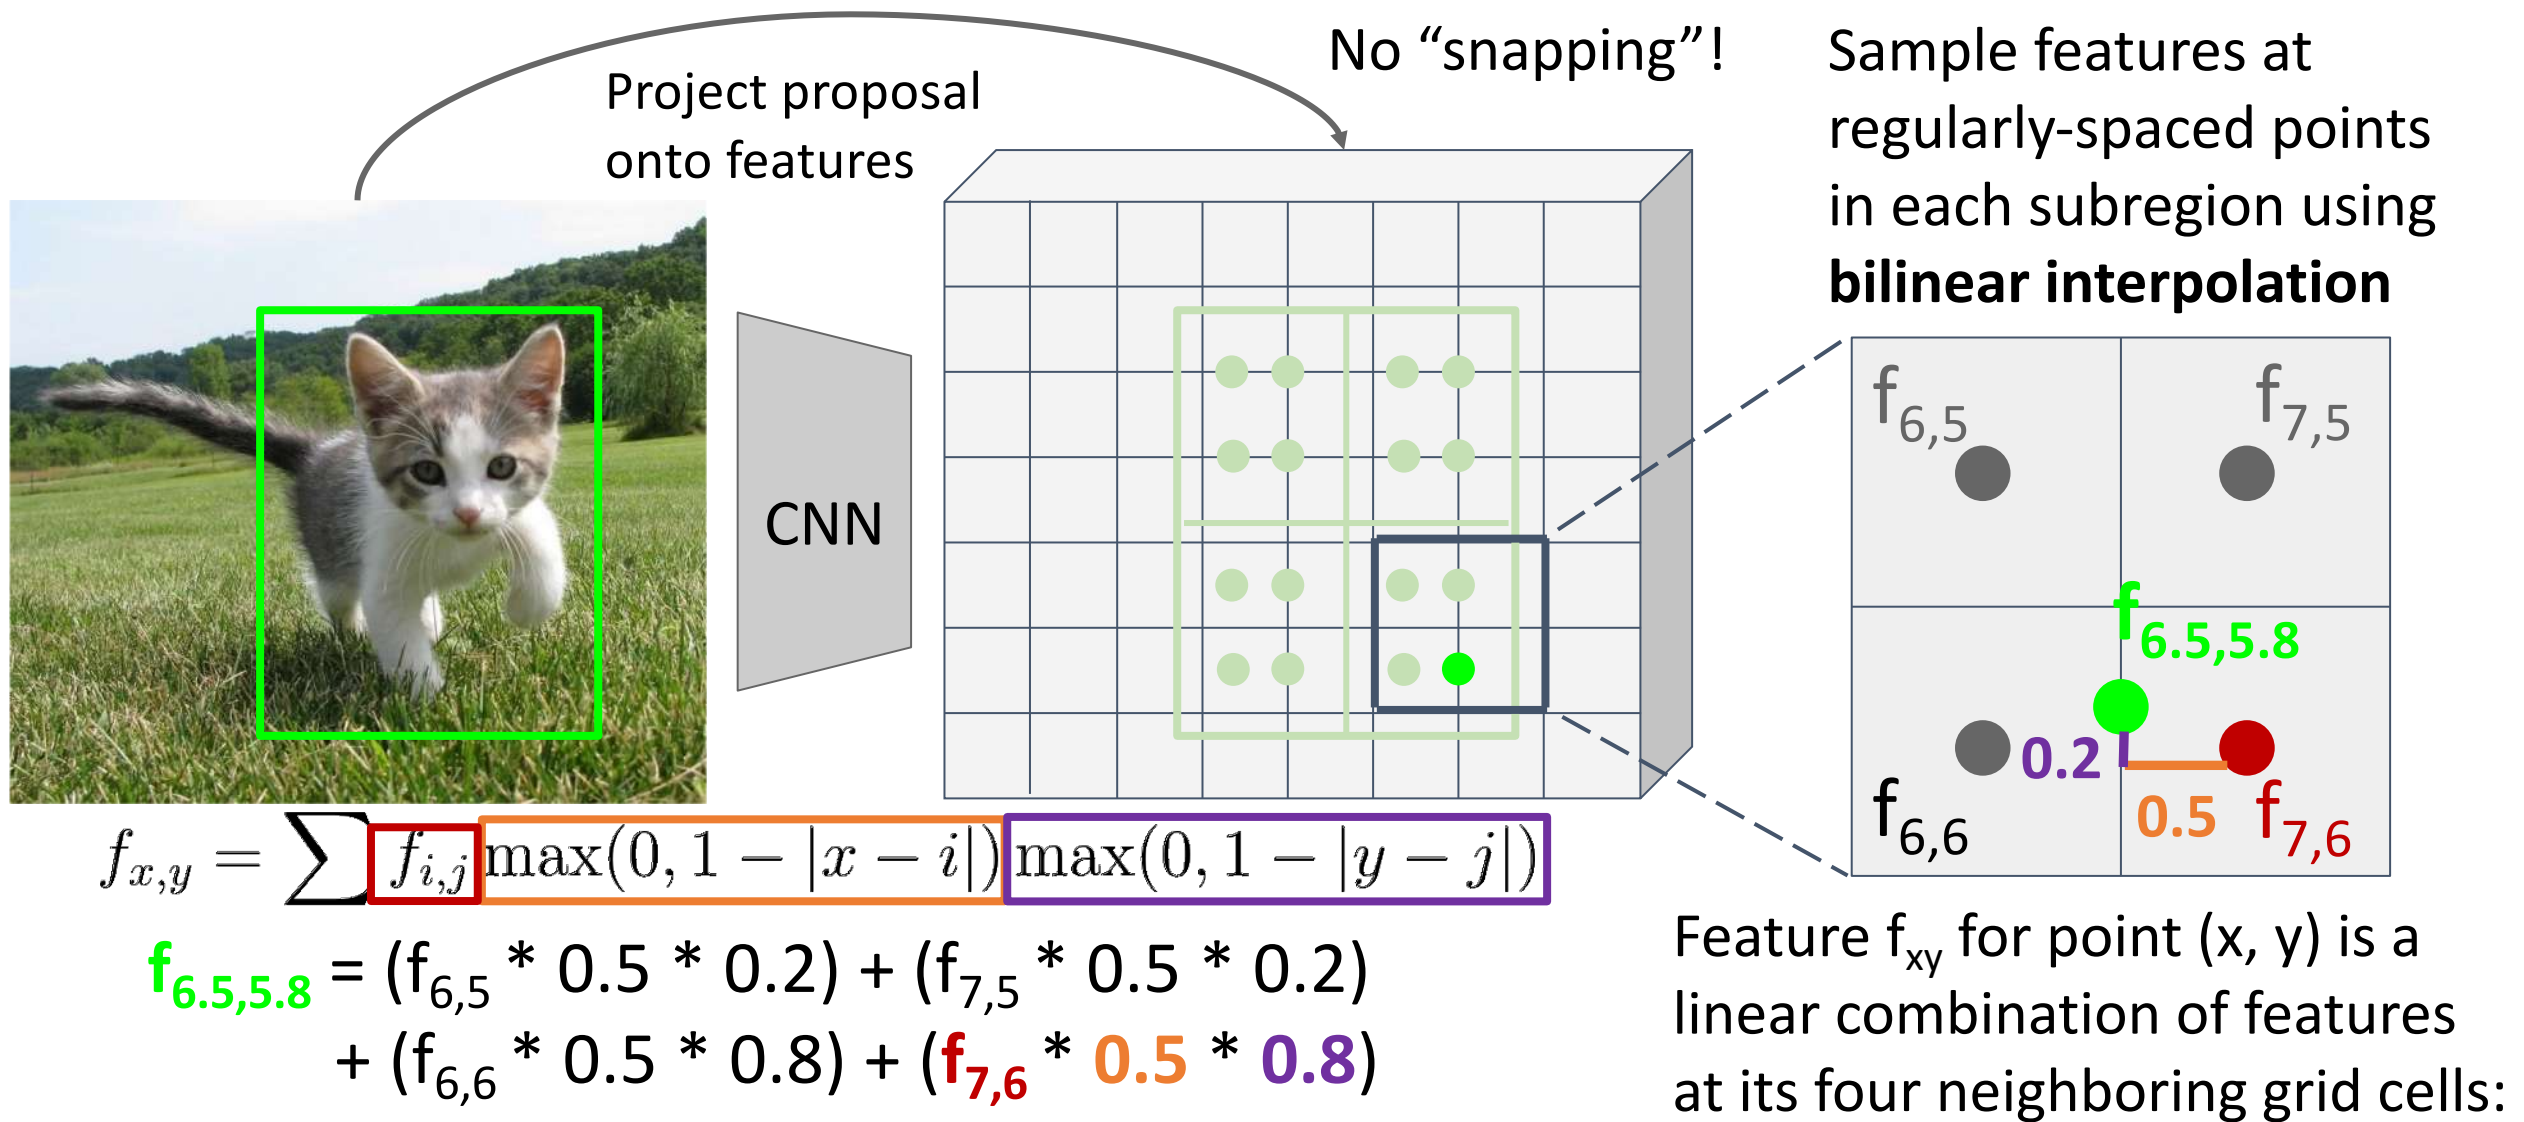


# Cropping Features: RoI Align

Divide into equal-sized subregions  
(may not be aligned to grid!)




# Cropping Features: RoI Align



# Cropping Features: RoI Align

Divide into equal-sized subregions  
(may not be aligned to grid!)



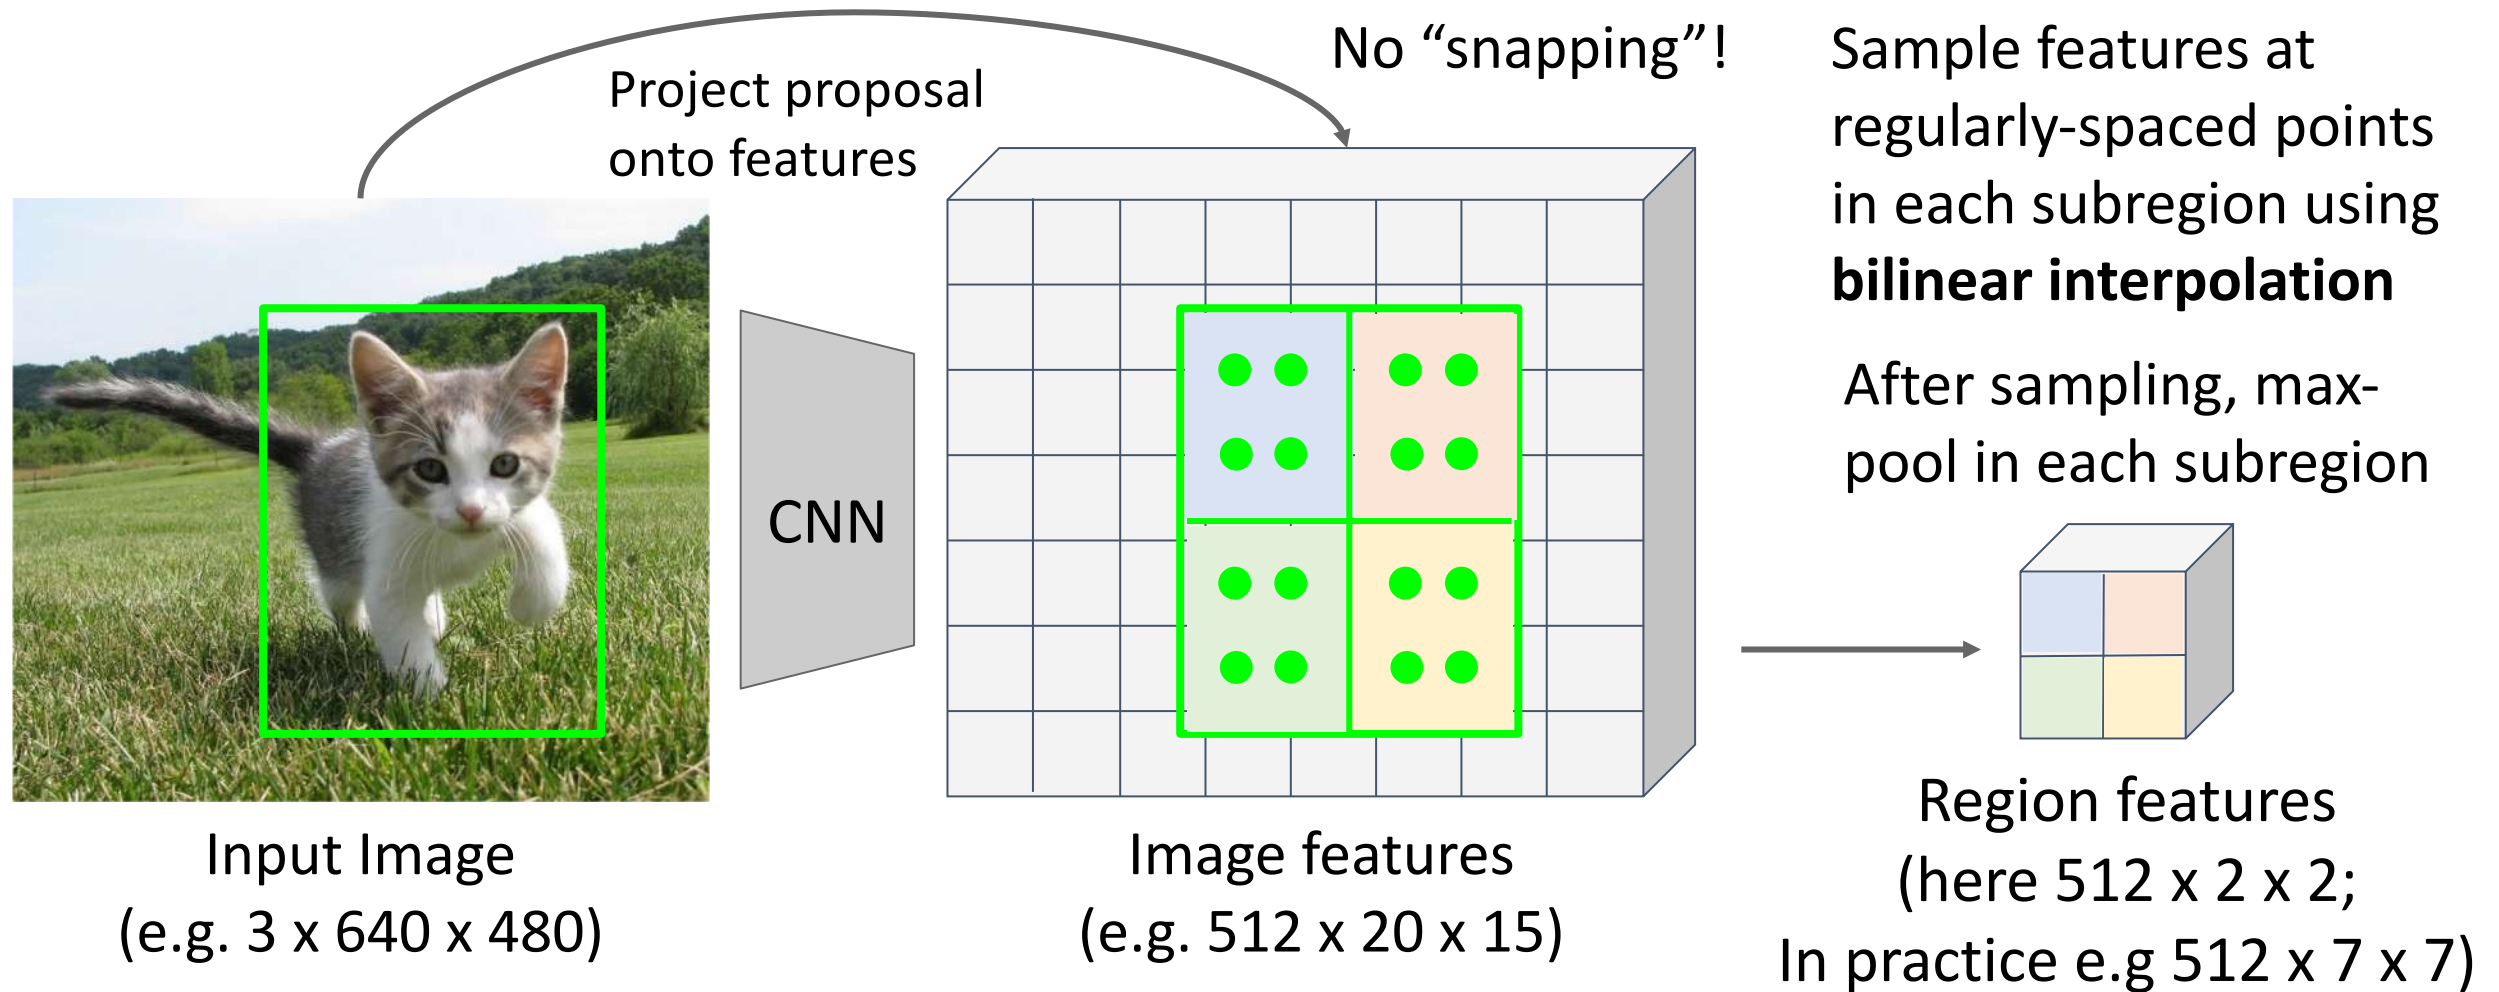
$$f_{x,y} = \sum_{i,j} f_{i,j} \max(0, 1 - |x - i|) \max(0, 1 - |y - j|)$$

This is differentiable! Upstream gradient for sampled feature will flow backward into each of the four nearest-neighbor gridpoints

Feature  $f_{xy}$  for point  $(x, y)$  is a linear combination of features at its four neighboring grid cells:

# Cropping Features: RoI Align

Divide into equal-sized subregions  
(may not be aligned to grid!)



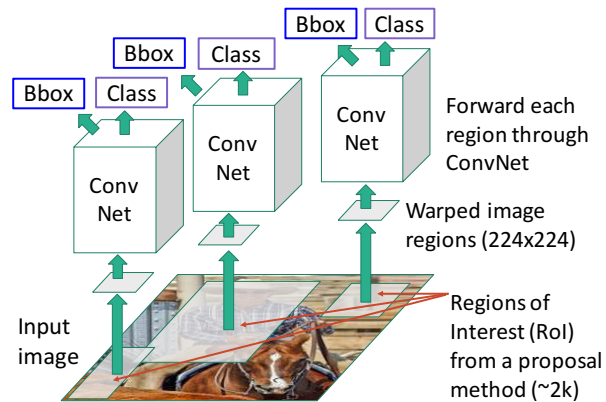
Output features now aligned to input box! And we can backprop to box coordinates!



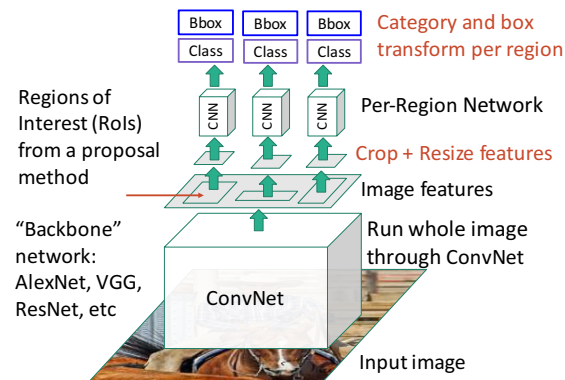
# Object Detection Methods

Both of these rely on anchor boxes.  
Can we do detection without anchors?


**“Slow” R-CNN:** Run CNN independently for each region




**Fast R-CNN:** Apply differentiable cropping to shared image features




**Faster R-CNN:**  
Compute proposals with CNN



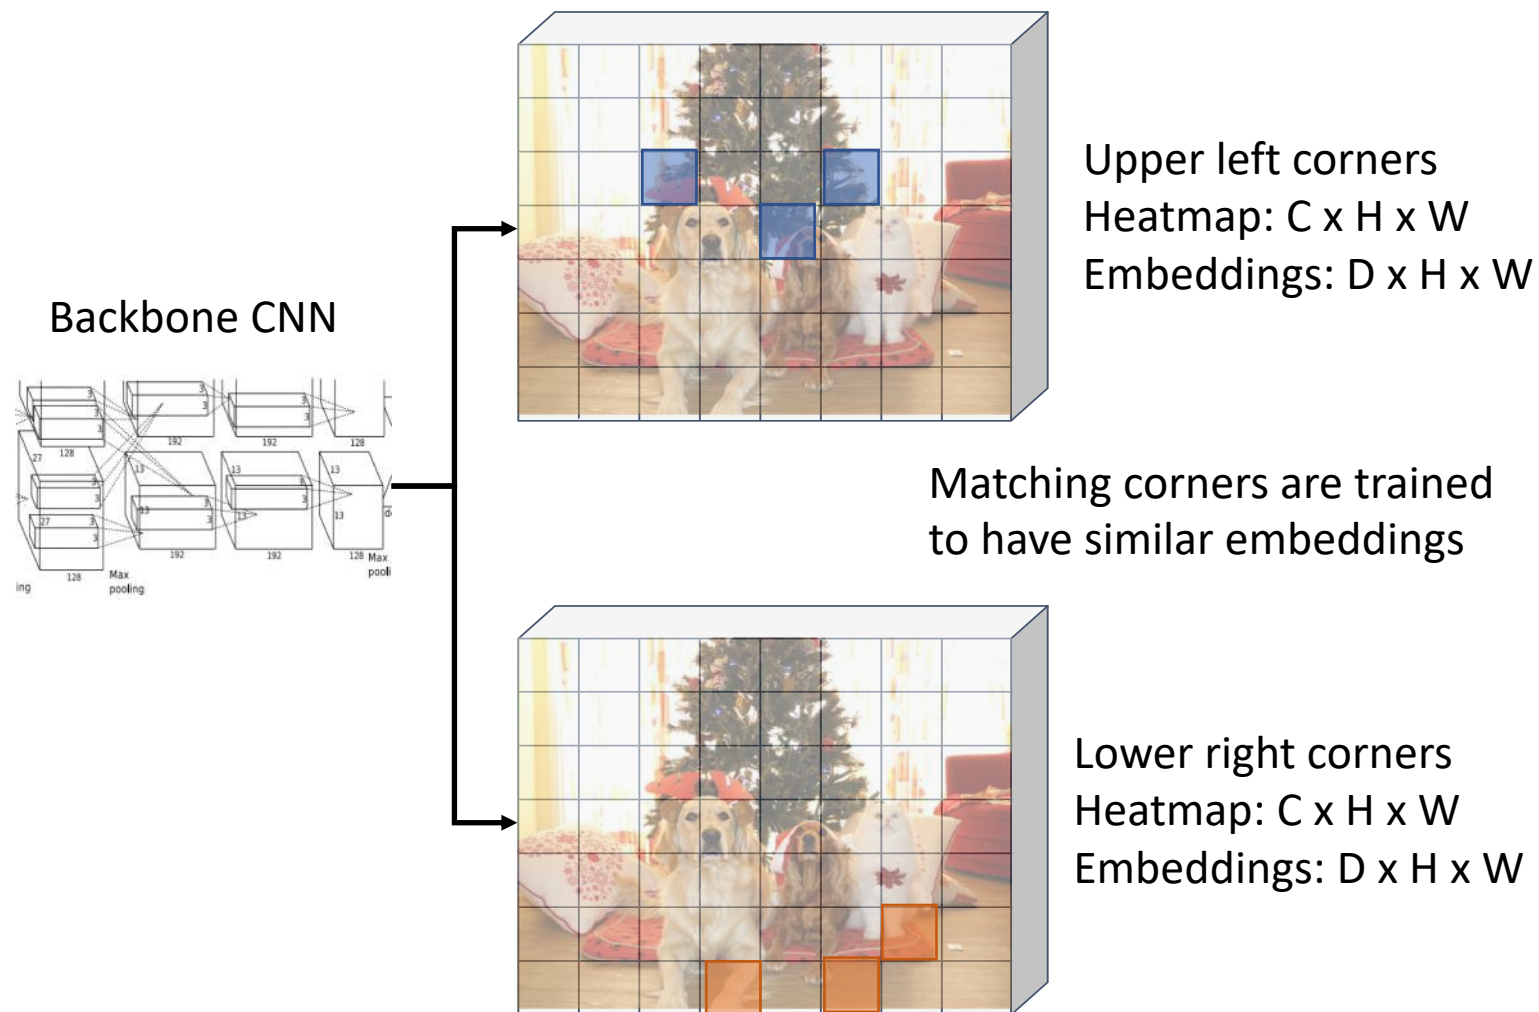
**Single-Stage:**  
Fully convolutional detector



# Detection without Anchors: CornerNet



## Represent bounding boxes by pairs of corners



Law and Deng, "CornerNet: Detecting Objects as Paired Keypoints", ECCV 2018



# Computer Vision Tasks: Object Detection


Classification



CAT

No spatial extent


Semantic Segmentation



GRASS, CAT, TREE,  
SKY

No objects, just pixels

Object Detection



DOG, DOG, CAT

Multiple Objects

Instance Segmentation



DOG, DOG, CAT

# Computer Vision Tasks: Semantic Segmentation


Classification



CAT

No spatial extent

Semantic Segmentation



GRASS, CAT, TREE,  
SKY

No objects, just pixels

Object Detection



DOG, DOG, CAT

Multiple Objects

Instance Segmentation




DOG, DOG, CAT


# Semantic Segmentation

Label each pixel in the image with a category label

Don't differentiate instances, only care about pixels

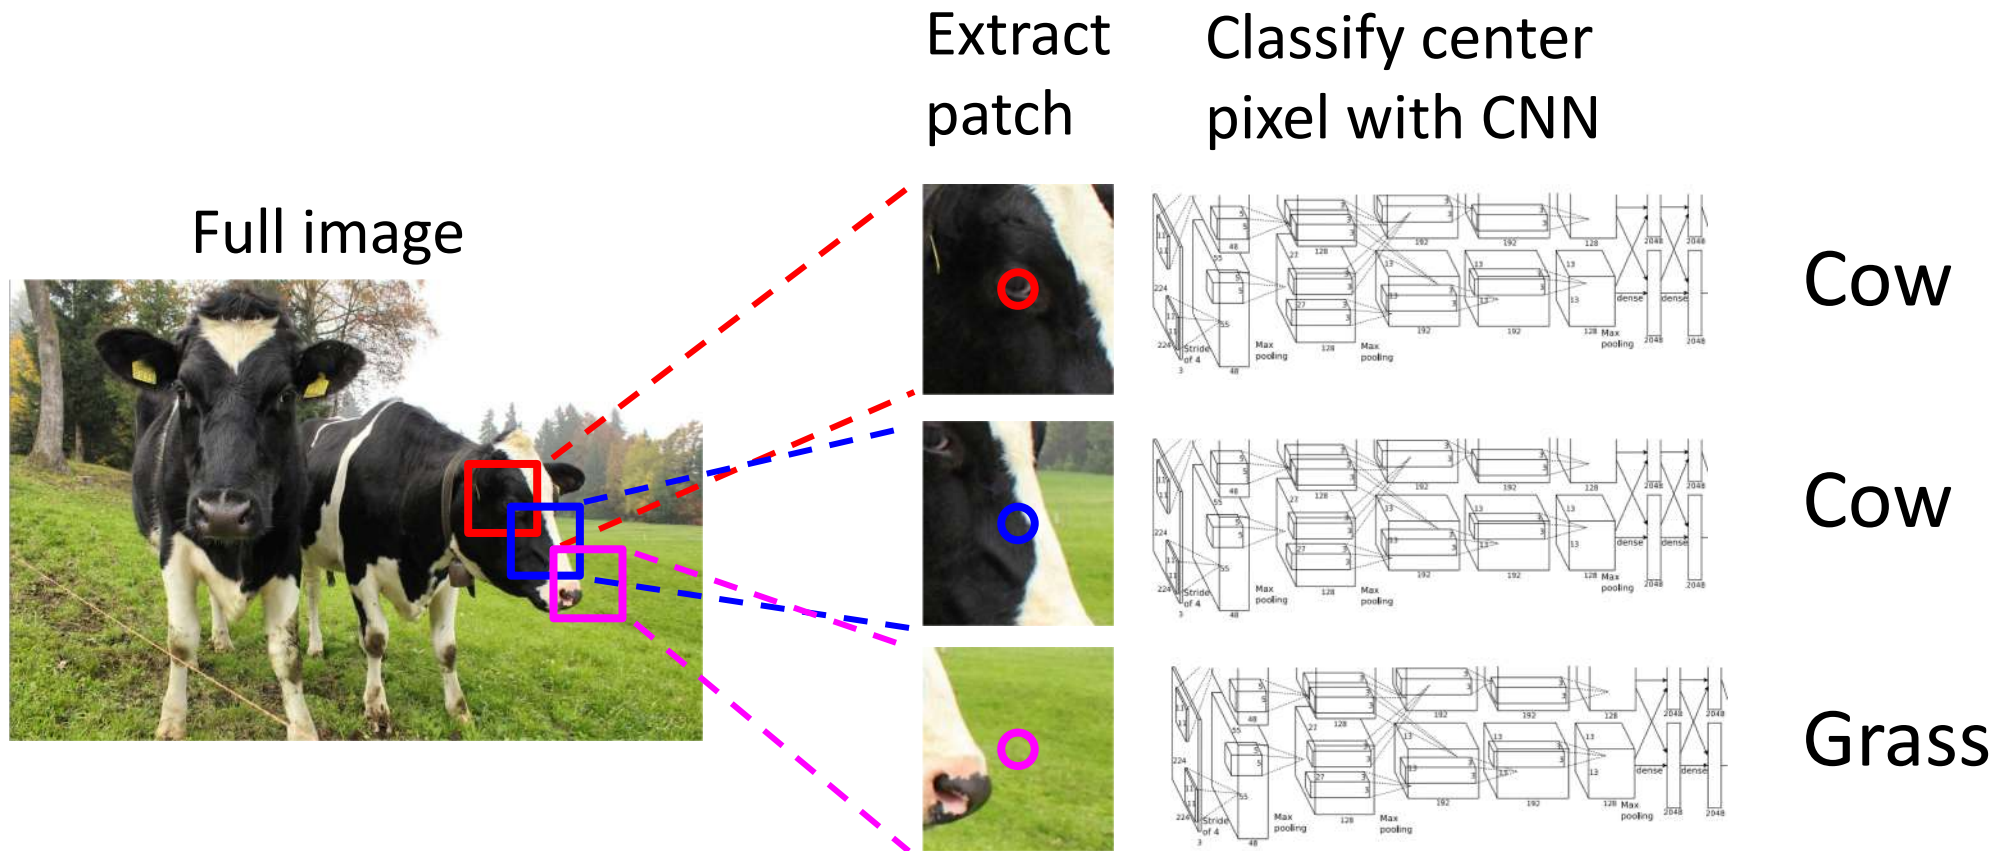


[This image](#) is [CC0 public domain](#)





# Semantic Segmentation Idea: Sliding Window



Farabet et al, "Learning Hierarchical Features for Scene Labeling," TPAMI 2013

Pinheiro and Collobert, "Recurrent Convolutional Neural Networks for Scene Labeling", ICML 2014





# Semantic Segmentation Idea: Sliding Window


## Extract patch

# Classify center pixel with CNN


## Full image



# Cow



# Cow



# Grass


Problem: Very inefficient! Not reusing shared features between overlapping patches

Farabet et al, "Learning Hierarchical Features for Scene Labeling," TPAMI 2013

Pinheiro and Collobert, "Recurrent Convolutional Neural Networks for Scene Labeling", ICML 2014

# Semantic Segmentation: Fully Convolutional Network

Design a network as a bunch of convolutional layers to make predictions for pixels all at once!




Loss function: Per-Pixel cross-entropy

Long et al, "Fully convolutional networks for semantic segmentation", CVPR 2015

# Semantic Segmentation: Fully Convolutional Network

Design a network as a bunch of convolutional layers to make predictions for pixels all at once!




Input:  $3 \times H \times W$

**Problem #1:** Effective receptive field size is linear in number of conv layers: With  $L$   $3 \times 3$  conv layers, receptive field is  $1 + 2L$

Long et al, "Fully convolutional networks for semantic segmentation", CVPR 2015

# Semantic Segmentation: Fully Convolutional Network

Design a network as a bunch of convolutional layers to make predictions for pixels all at once!



Input:  
 $3 \times H \times W$

**Problem #1:** Effective receptive field size is linear in number of conv layers: With  $L$   $3 \times 3$  conv layers, receptive field is  $1+2L$


**Problem #2:** Convolution on high res images is expensive! Recall ResNet stem aggressively downsamples

Long et al, "Fully convolutional networks for semantic segmentation", CVPR 2015



# Semantic Segmentation: Fully Convolutional Network

Design network as a bunch of convolutional layers, with **downsampling** and **upsampling** inside the network!



# Semantic Segmentation: Fully Convolutional Network


**Downsampling:**  
Pooling, strided  
convolution

Design network as a bunch of convolutional layers, with **downsampling** and **upsampling** inside the network!

**Upsampling:**  
???




Input:  
 $3 \times H \times W$



Predictions:  
 $H \times W$

# In-Network Upsampling: “Unpooling”

## Bed of Nails




Input  
 $C \times 2 \times 2$

Output  
 $C \times 4 \times 4$

# In-Network Upsampling: “Unpooling”


## Bed of Nails



Input  
 $C \times 2 \times 2$

Output  
 $C \times 4 \times 4$

## Nearest Neighbor




Input  
 $C \times 2 \times 2$

Output  
 $C \times 4 \times 4$



# In-Network Upsampling: Bilinear Interpolation



1.00	1.25	1.75	2.00
1.50	1.75	2.25	2.50
2.50	2.75	3.25	3.50
3.00	3.25	3.75	4.00


Input: C x 2 x 2

Output: C x 4 x 4

$$f_{x,y} = \sum_{i,j} f_{i,j} \max(0, 1 - |x - i|) \max(0, 1 - |y - j|) \quad \begin{array}{l} i \in \{\lfloor x \rfloor - 1, \dots, \lceil x \rceil + 1\} \\ j \in \{\lfloor y \rfloor - 1, \dots, \lceil y \rceil + 1\} \end{array}$$

Use two closest neighbors in x and y  
to construct linear approximations

# In-Network Upsampling: Bicubic Interpolation



Input:  $C \times 2 \times 2$

Output:  $C \times 4 \times 4$

Use **three** closest neighbors in x and y to construct **cubic** approximations  
(This is how we normally resize images!)

# In-Network Upsampling: “Max Unpooling”

**Max Pooling:** Remember which position had the max

1	2	6	3
3	5	2	1
1	2	2	1
7	3	4	8



5	6
7	8




Rest  
of  
net



1	2
3	4



0	0	2	0
0	1	0	0
0	0	0	0
3	0	0	4




Pair each downsampling layer  
with an upsampling layer


Noh et al, “Learning Deconvolution Network for Semantic Segmentation”, ICCV 2015

# Learnable Upsampling: Transposed Convolution

**Recall:** Normal 3 x 3 convolution, stride 1, pad 1



Input: 4 x 4

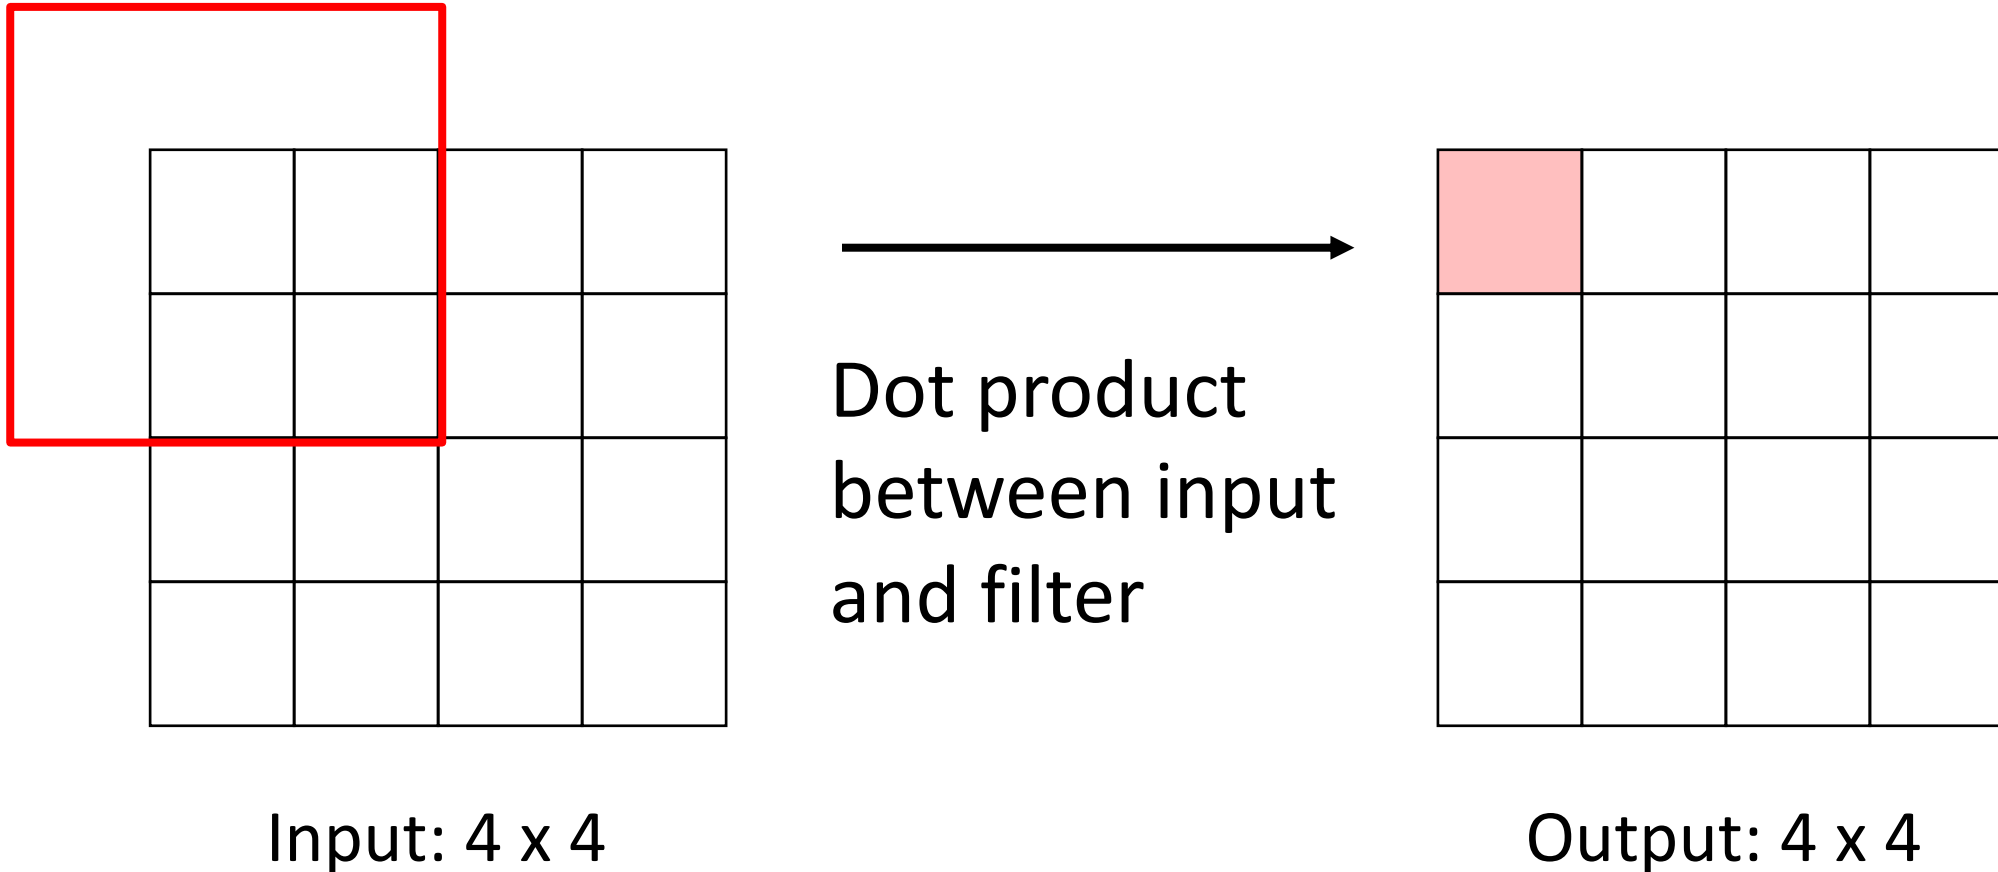


Output: 4 x 4




# Learnable Upsampling: Transposed Convolution

**Recall:** Normal 3 x 3 convolution, stride 1, pad 1



# Learnable Upsampling: Transposed Convolution


**Recall:** Normal 3 x 3 convolution, stride 1, pad 1



Input: 4 x 4




Dot product  
between input  
and filter




Output: 4 x 4

# Learnable Upsampling: Transposed Convolution

**Recall:** Normal 3 x 3 convolution, stride 2, pad 1




Input: 4 x 4



Output: 2 x 2

# Learnable Upsampling: Transposed Convolution


**Recall:** Normal 3 x 3 convolution, stride 2, pad 1



Input: 4 x 4



Dot product  
between input  
and filter




Output: 2 x 2



# Learnable Upsampling: Transposed Convolution


**Recall:** Normal 3 x 3 convolution, stride 2, pad 1



Input: 4 x 4




Dot product  
between input  
and filter



Output: 2 x 2

# Learnable Upsampling: Transposed Convolution

**Recall:** Normal 3 x 3 convolution, stride 2, pad 1




Input: 4 x 4

Convolution with stride  $> 1$  is “Learnable Downsampling”  
Can we use stride  $< 1$  for “Learnable Upsampling”?




Dot product  
between input  
and filter




Output: 2 x 2

# Learnable Upsampling: Transposed Convolution

**3 x 3 convolution transpose, stride 2**




Input: 2 x 2



Output: 4 x 4

# Learnable Upsampling: Transposed Convolution


3 x 3 **convolution transpose**, stride 2



Input: 2 x 2



Weight filter by  
input value and  
copy to output




Output: 4 x 4

# Learnable Upsampling: Transposed Convolution

3 x 3 **convolution transpose**, stride 2


Filter moves 2 pixels in output  
for every 1 pixel in input



Input: 2 x 2



Weight filter by  
input value and  
copy to output




Output: 4 x 4



# Learnable Upsampling: Transposed Convolution

3 x 3 **convolution transpose**, stride 2


Filter moves 2 pixels in output  
for every 1 pixel in input



Input: 2 x 2



Weight filter by  
input value and  
copy to output




Output: 4 x 4

# Learnable Upsampling: Transposed Convolution

3 x 3 **convolution transpose**, stride 2


This gives 5x5 output – need to trim one pixel from top and left to give 4x4 output



Input: 2 x 2




Weight filter by  
input value and  
copy to output




Output: 4 x 4

# Transposed Convolution: 1D example


**Input**



**Filter**



**Output**




Output has copies of filter weighted by input

Stride 2: Move 2 pixels output for each pixel in input


Sum at overlaps

# Transposed Convolution: 1D example


**Input**



**Filter**



**Output**



This has many names:

- Deconvolution (bad)!
- Upconvolution
- Fractionally strided convolution
- Backward strided convolution
- Transposed Convolution (best name)

# Convolution as Matrix Multiplication (1D Example)

We can express convolution in terms of a matrix multiplication

$$\vec{x} * \vec{a} = X \vec{a}$$

$$\begin{bmatrix} x & y & x & 0 & 0 & 0 \\ 0 & x & y & x & 0 & 0 \\ 0 & 0 & x & y & x & 0 \\ 0 & 0 & 0 & x & y & x \end{bmatrix} \begin{bmatrix} 0 \\ a \\ b \\ c \\ d \\ 0 \end{bmatrix} = \begin{bmatrix} ay + bz \\ ax + by + cz \\ bx + cy + dz \\ cx + dy \end{bmatrix}$$

Example: 1D conv, kernel  
size=3, stride=1, padding=1



# Convolution as Matrix Multiplication (1D Example)

We can express convolution in terms of a matrix multiplication

$$\vec{x} * \vec{a} = X \vec{a}$$

$$\begin{bmatrix} x & y & x & 0 & 0 & 0 \\ 0 & x & y & x & 0 & 0 \\ 0 & 0 & x & y & x & 0 \\ 0 & 0 & 0 & x & y & x \end{bmatrix} \begin{bmatrix} 0 \\ a \\ b \\ c \\ d \\ 0 \end{bmatrix} = \begin{bmatrix} ay + bz \\ ax + by + cz \\ bx + cy + dz \\ cx + dy \end{bmatrix}$$

Example: 1D conv, kernel size=3, stride=1, padding=1

**Transposed convolution** multiplies by the transpose of the same matrix:

$$\vec{x} *^T \vec{a} = X^T \vec{a}$$

$$\begin{bmatrix} x & 0 & 0 & 0 \\ y & x & 0 & 0 \\ z & y & x & 0 \\ 0 & z & y & x \\ 0 & 0 & z & y \\ 0 & 0 & 0 & z \end{bmatrix} \begin{bmatrix} a \\ b \\ c \\ d \end{bmatrix} = \begin{bmatrix} ax \\ ay + bx \\ az + by + cx \\ bz + cy + dx \\ cz + dy \\ dz \end{bmatrix}$$

When stride=1, transposed conv is just a regular conv (with different padding rules)

# Convolution as Matrix Multiplication (1D Example)

We can express convolution in terms of a matrix multiplication

$$\vec{x} * \vec{a} = X \vec{a}$$

$$\begin{bmatrix} x & y & x & 0 & 0 & 0 \\ 0 & 0 & x & y & x & 0 \end{bmatrix} \begin{bmatrix} 0 \\ a \\ b \\ c \\ d \\ 0 \end{bmatrix} = \begin{bmatrix} ay + bz \\ bx + cy + dz \end{bmatrix}$$

Example: 1D conv, kernel  
size=3, stride=2, padding=1

**Transposed convolution** multiplies by the transpose of the same matrix:

$$\vec{x} *^T \vec{a} = X^T \vec{a}$$

# Convolution as Matrix Multiplication (1D Example)

We can express convolution in terms of a matrix multiplication

$$\vec{x} * \vec{a} = X \vec{a}$$

$$\begin{bmatrix} x & y & x & 0 & 0 & 0 \\ 0 & 0 & x & y & x & 0 \end{bmatrix} \begin{bmatrix} 0 \\ a \\ b \\ c \\ d \\ 0 \end{bmatrix} = \begin{bmatrix} ay + bz \\ bx + cy + dz \end{bmatrix}$$

Example: 1D conv, kernel size=3, stride=2, padding=1

**Transposed convolution** multiplies by the transpose of the same matrix:

$$\vec{x} *^T \vec{a} = X^T \vec{a}$$

$$\begin{bmatrix} x & 0 \\ y & 0 \\ z & x \\ 0 & y \\ 0 & z \\ 0 & 0 \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} = \begin{bmatrix} ax \\ ay \\ az + bx \\ by \\ bz \\ 0 \end{bmatrix}$$

When stride>1, transposed convolution cannot be expressed as normal conv

# Semantic Segmentation: Fully Convolutional Network


**Downsampling:**  
Pooling, strided  
convolution

Design network as a bunch of convolutional layers, with **downsampling** and **upsampling** inside the network!

**Upsampling:**  
Interpolation,  
transposed conv



Input:  
 $3 \times H \times W$



Predictions:  
 $H \times W$


Loss function: Per-Pixel cross-entropy

# Computer Vision Tasks

**Object Detection:** Detects individual object instances, but only gives box



**Semantic Segmentation:** Gives per-pixel labels, but merges instances







# Things and Stuff

**Things:** Object categories that can be separated into object instances (e.g. cats, cars, person)



**Stuff:** Object categories that cannot be separated into instances (e.g. sky, grass, water, trees)




# Computer Vision Tasks

**Object Detection:** Detects individual object instances, but only gives box (Only things!)



**Semantic Segmentation:** Gives per-pixel labels, but merges instances (Both things and stuff)



# Computer Vision Tasks: Instance Segmentation


Classification



CAT

No spatial extent

Semantic Segmentation



GRASS, CAT, TREE,  
SKY

No objects, just pixels

Object Detection



DOG, DOG, CAT

Multiple Objects

Instance Segmentation




DOG, DOG, CAT



# Computer Vision Tasks: Instance Segmentation

## **Instance Segmentation:**

Detect all objects in the image, and identify the pixels that belong to each object (Only things!)




[This image](#) is [CC0 public domain](#)

# Computer Vision Tasks: Instance Segmentation

## **Instance Segmentation:**

Detect all objects in the image, and identify the pixels that belong to each object (Only things!)

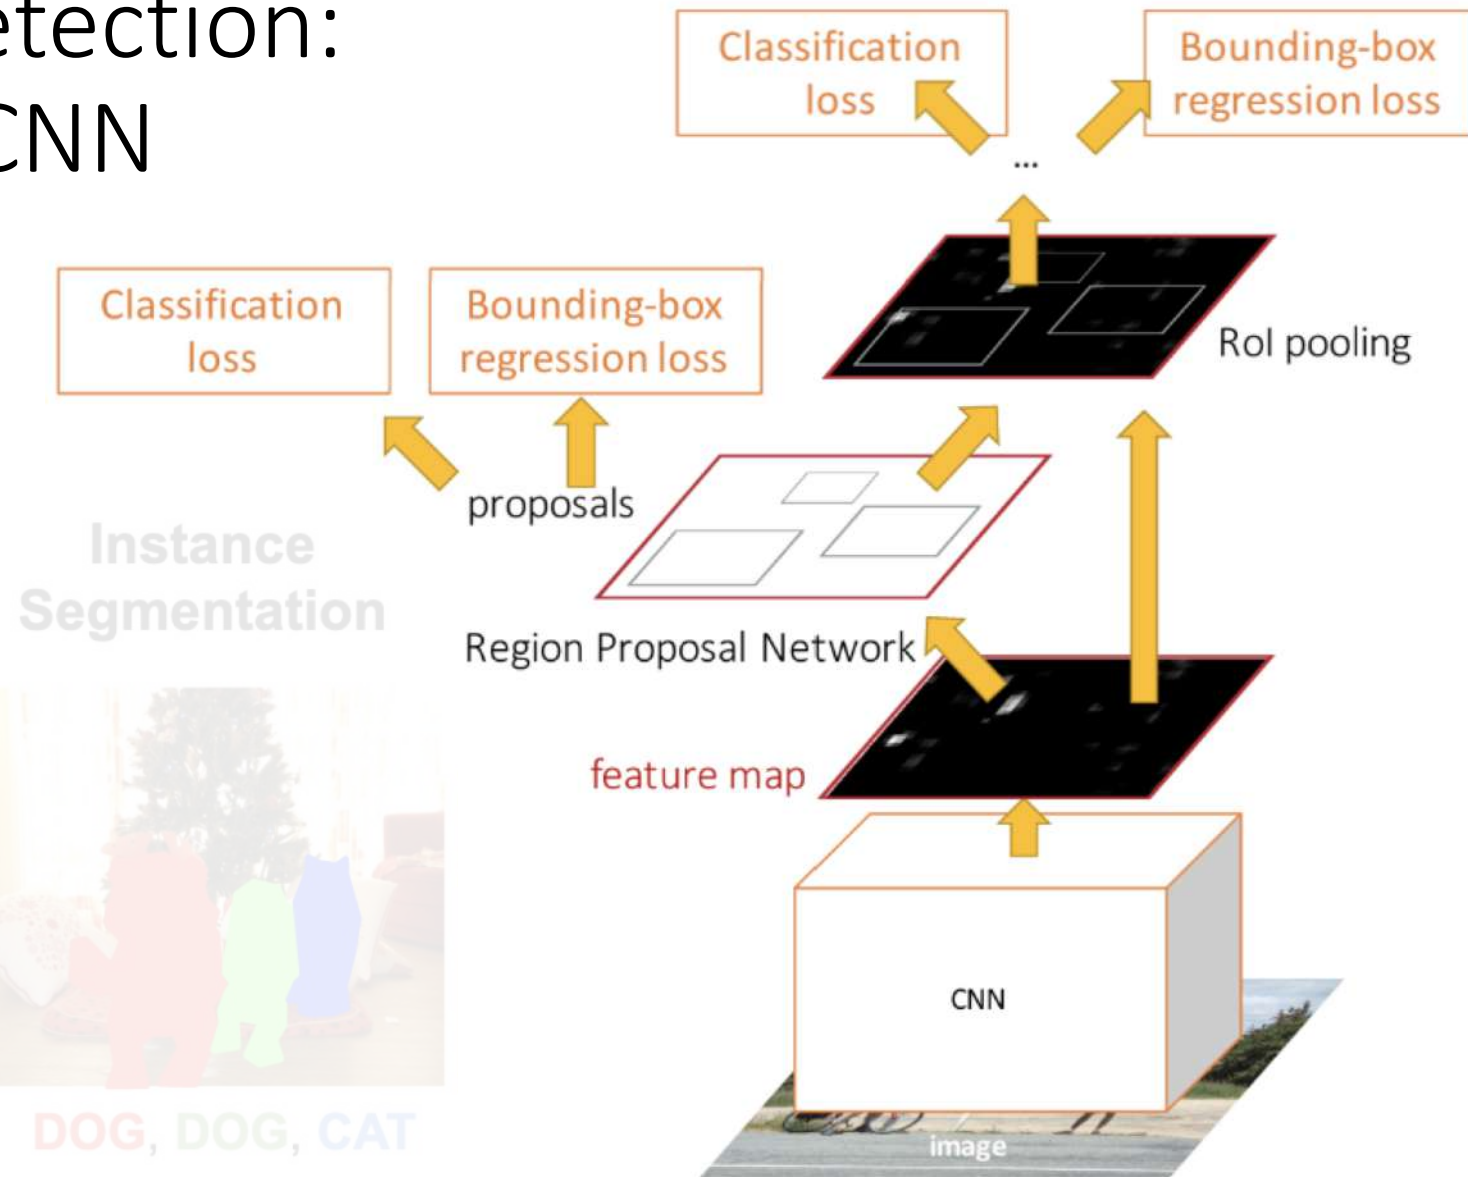
**Approach:** Perform object detection, then predict a segmentation mask for each object!



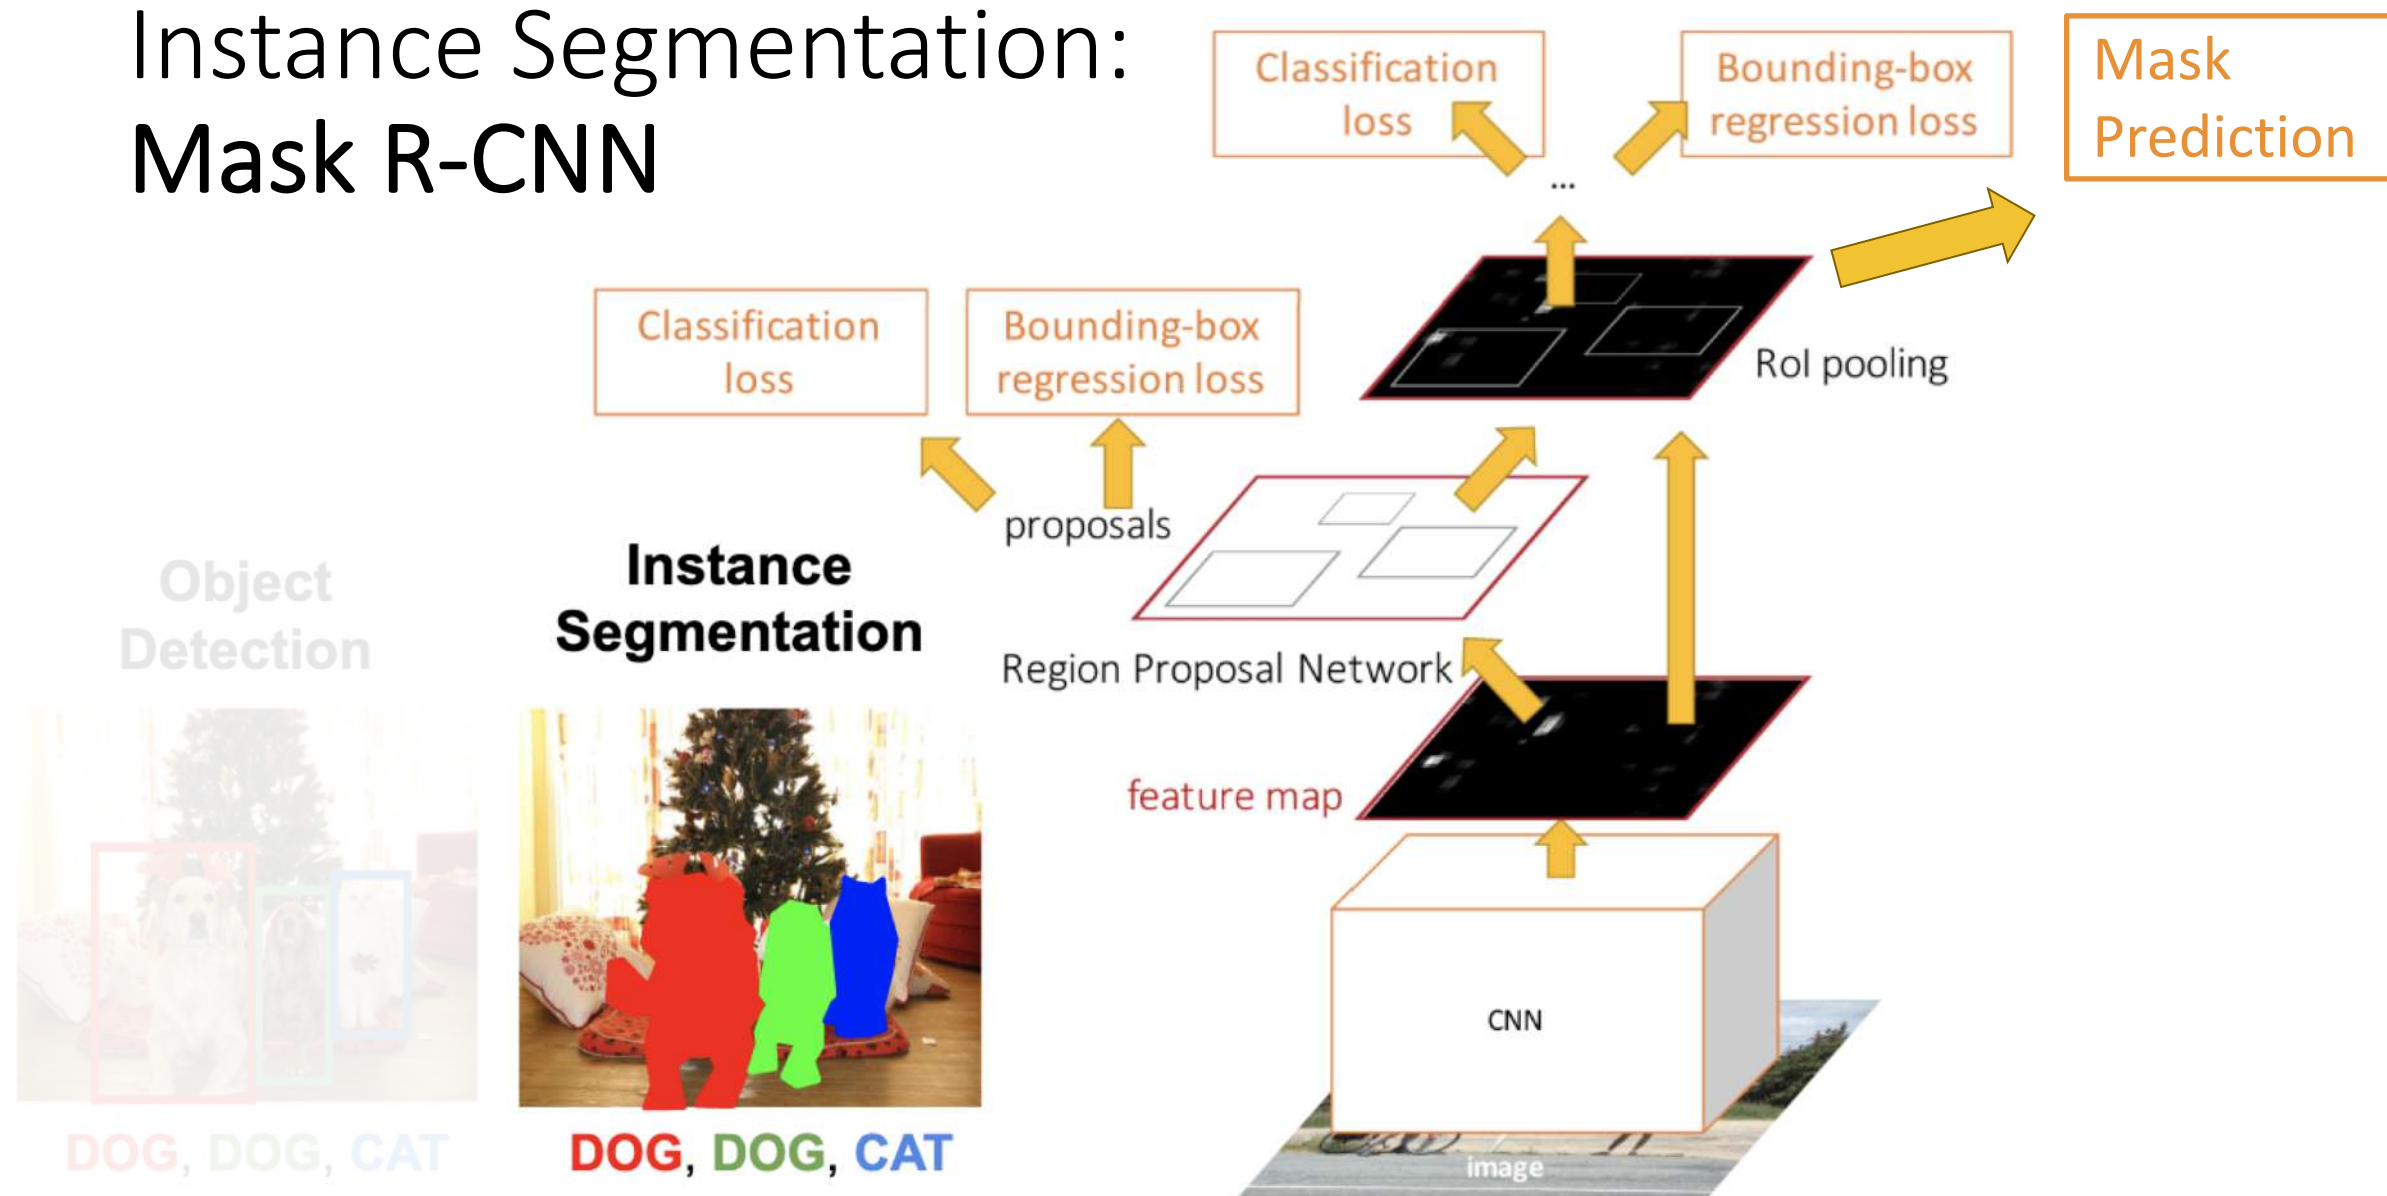
[This image](#) is [CC0 public domain](#)



# Object Detection: Faster R-CNN




# Instance Segmentation: Mask R-CNN




He et al, "Mask R-CNN", ICCV 2017

# Mask R-CNN




CNN  
+RPN






Predict a mask for  
each of  $C$  classes:  
 $C \times 28 \times 28$

# Mask R-CNN: Example Training Targets







# Mask R-CNN: Example Training Targets










# Mask R-CNN: Example Training Targets








# Mask R-CNN: Example Training Targets




# Mask R-CNN: Very Good Results!






# Beyond Instance Segmentation

**Instance Segmentation:** Separate object instances, but only things




**Semantic Segmentation:** Identify both things and stuff, but doesn't separate instances



# Beyond Instance Segmentation: Panoptic Segmentation

Label all pixels in the image (both things and stuff)


For “thing” categories also separate into instances



Kirillov et al, “Panoptic Segmentation”, CVPR 2019

Kirillov et al, “Panoptic Feature Pyramid Networks”, CVPR 2019

# Beyond Instance Segmentation: Panoptic Segmentation



Kirillov et al, "Panoptic Feature Pyramid Networks", CVPR 2019



# Beyond Instance Segmentation: Human Keypoints

Represent the pose of a human by locating a set of **keypoints**

e.g. 17 keypoints:


- Nose
- Left / Right eye
- Left / Right ear
- Left / Right shoulder
- Left / Right elbow
- Left / Right wrist
- Left / Right hip
- Left / Right knee
- Left / Right ankle



[Person image](#) is [CC0 public domain](#)



# Mask R-CNN: Instance Segmentation



Object  
Detection

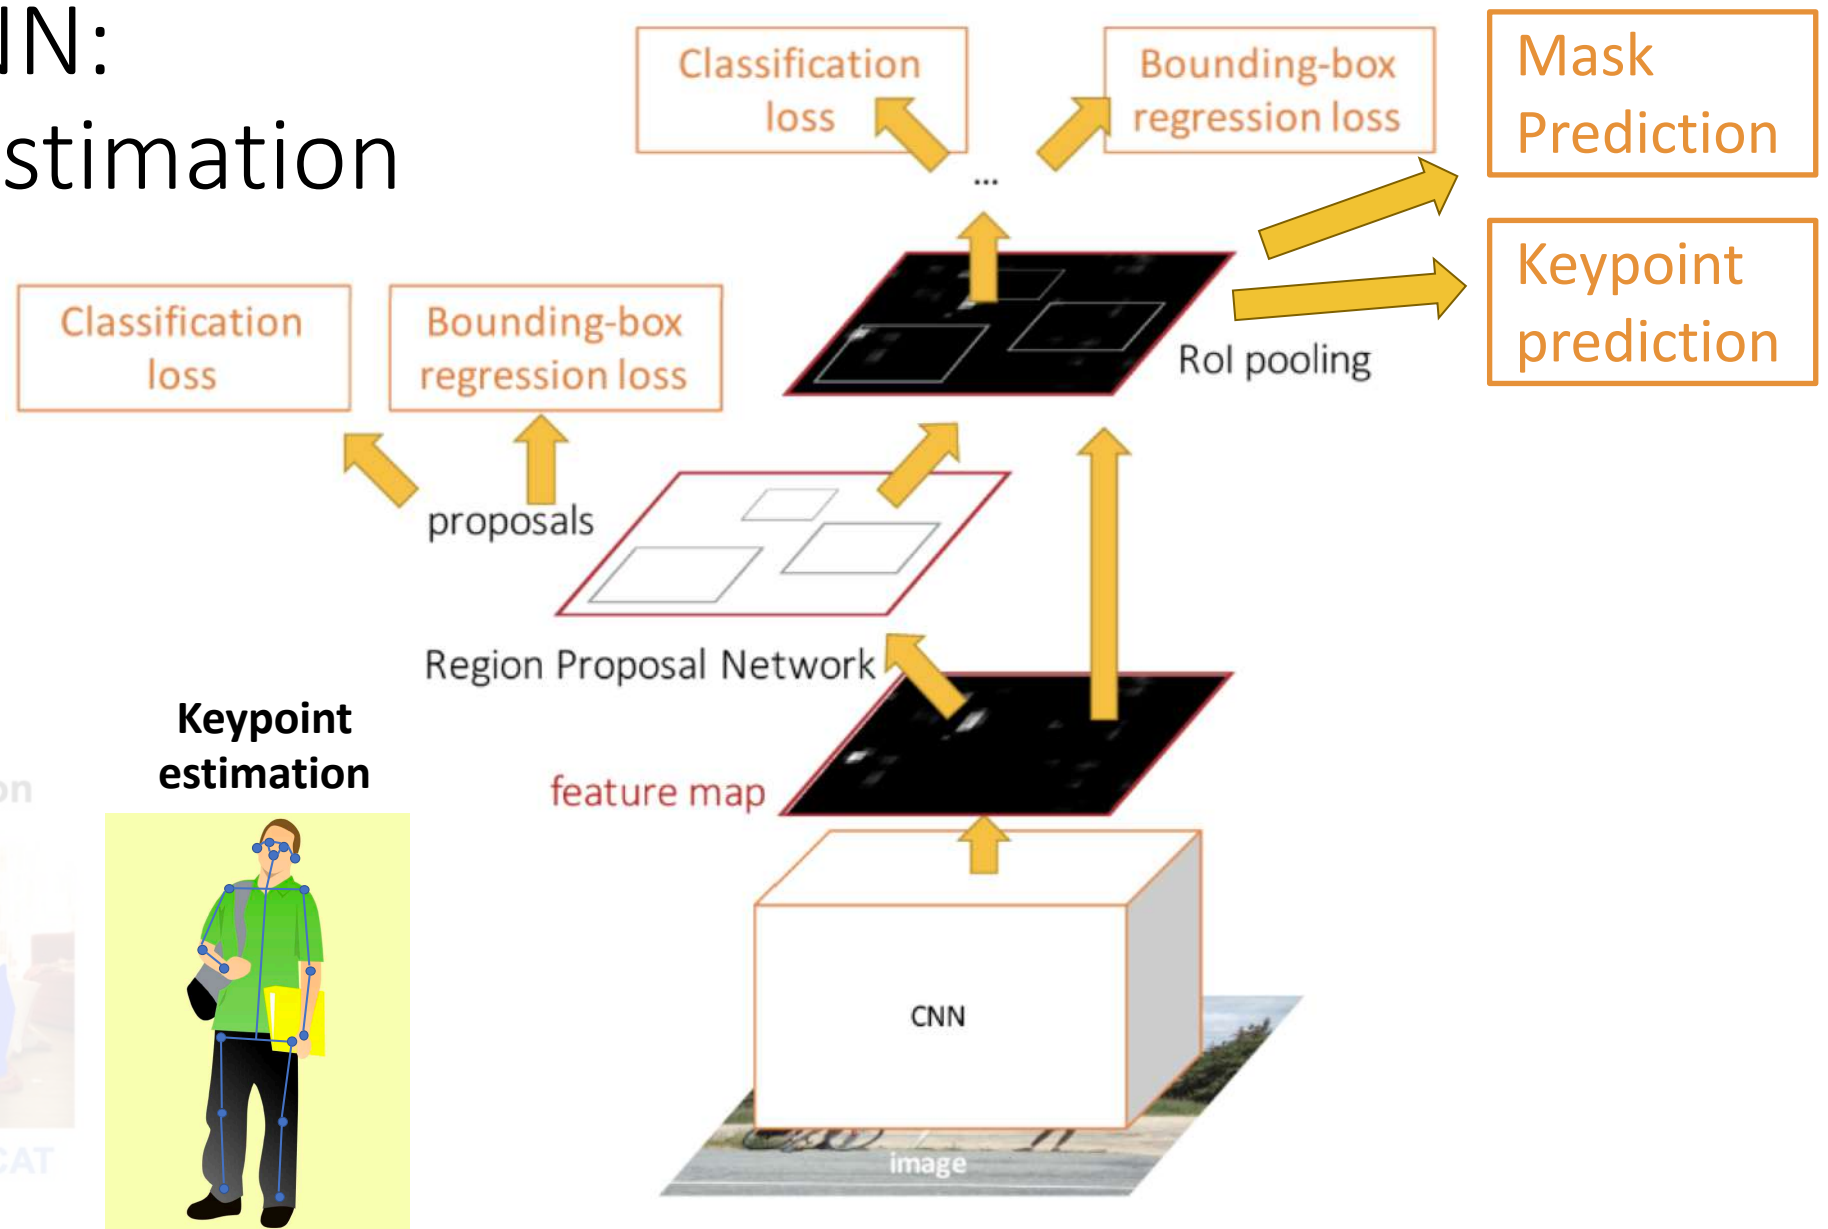
Instance  
Segmentation

Keypoint  
estimation



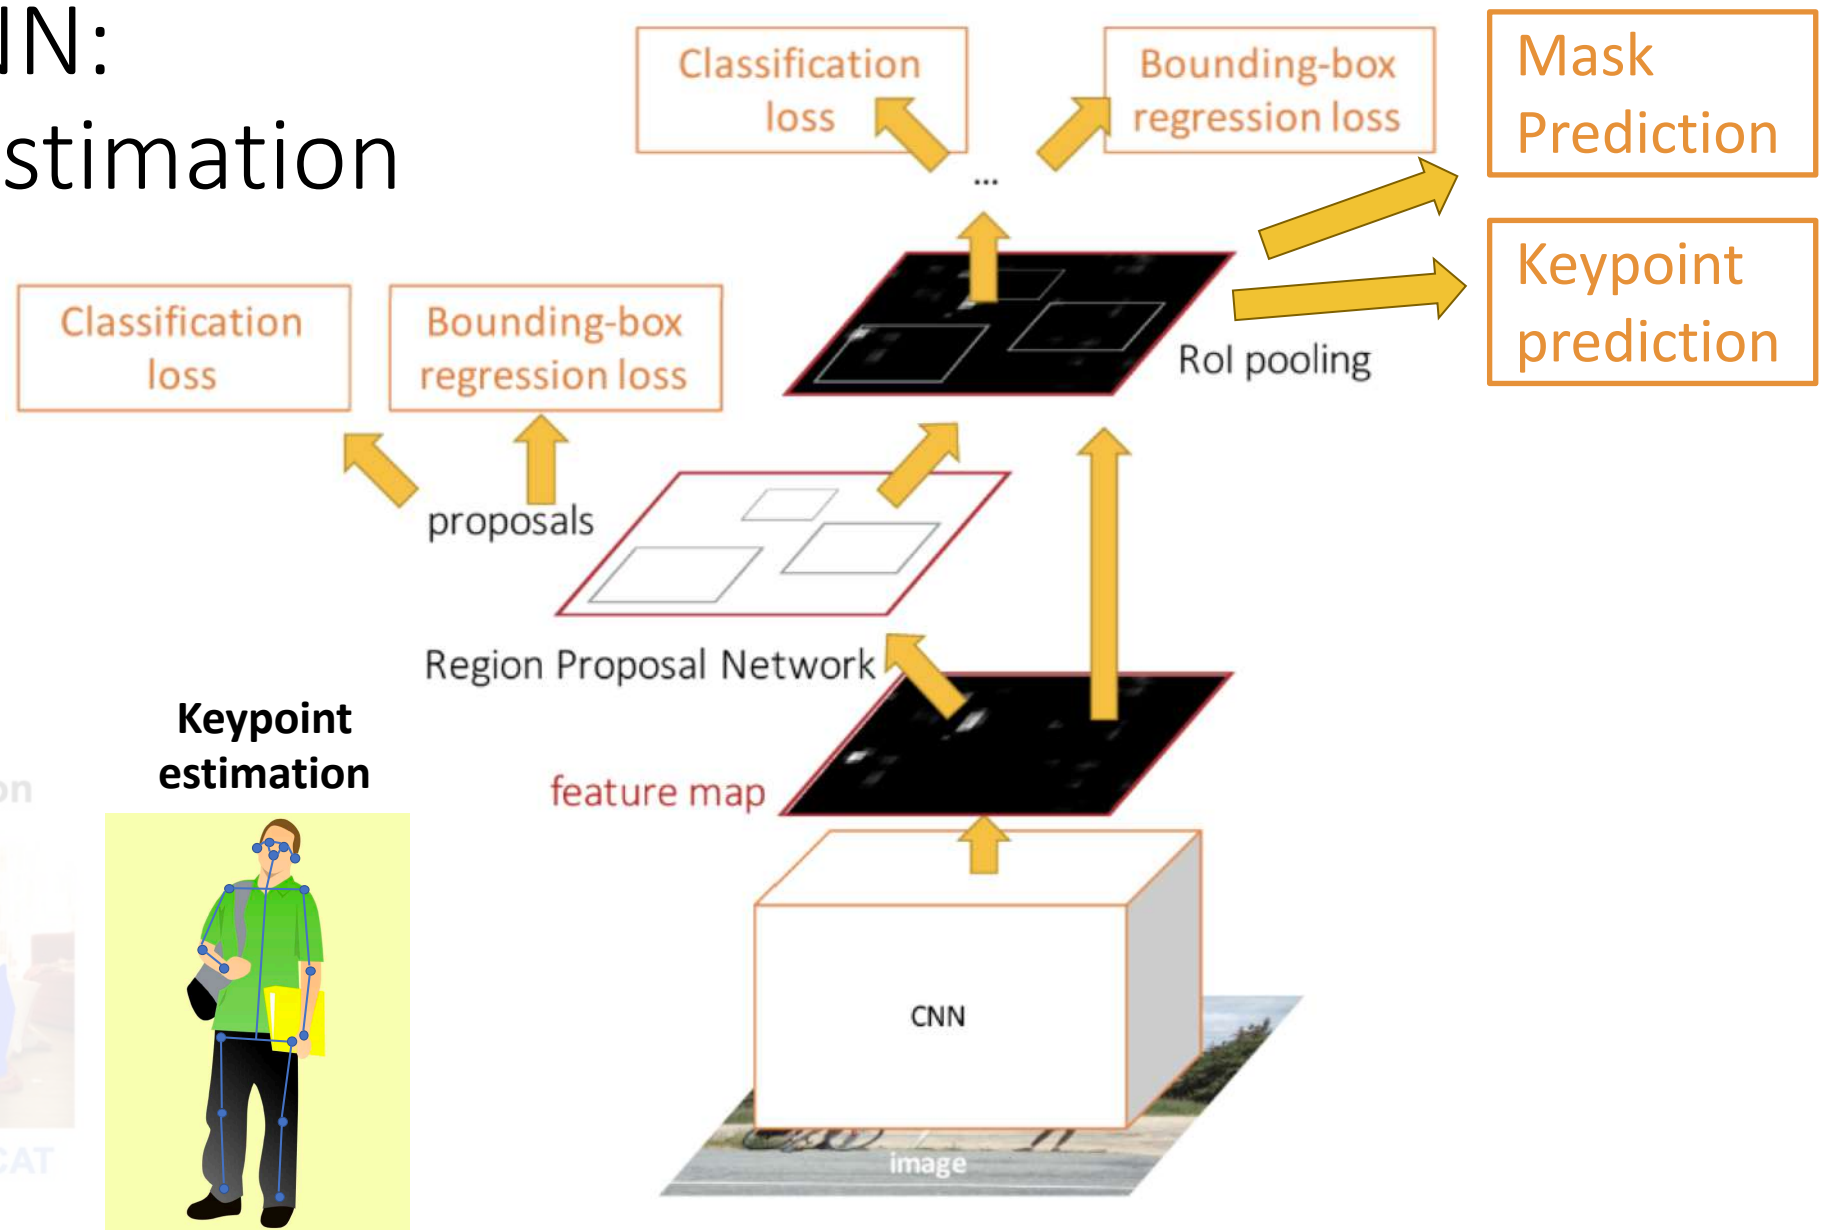
He et al, "Mask R-CNN", ICCV 2017

# Mask R-CNN: Keypoint Estimation




He et al, "Mask R-CNN", ICCV 2017

# Mask R-CNN: Keypoint Estimation



He et al, "Mask R-CNN", ICCV 2017


# Mask R-CNN: Keypoints



Ground-truth has one "pixel" turned on per keypoint. Train with softmax loss

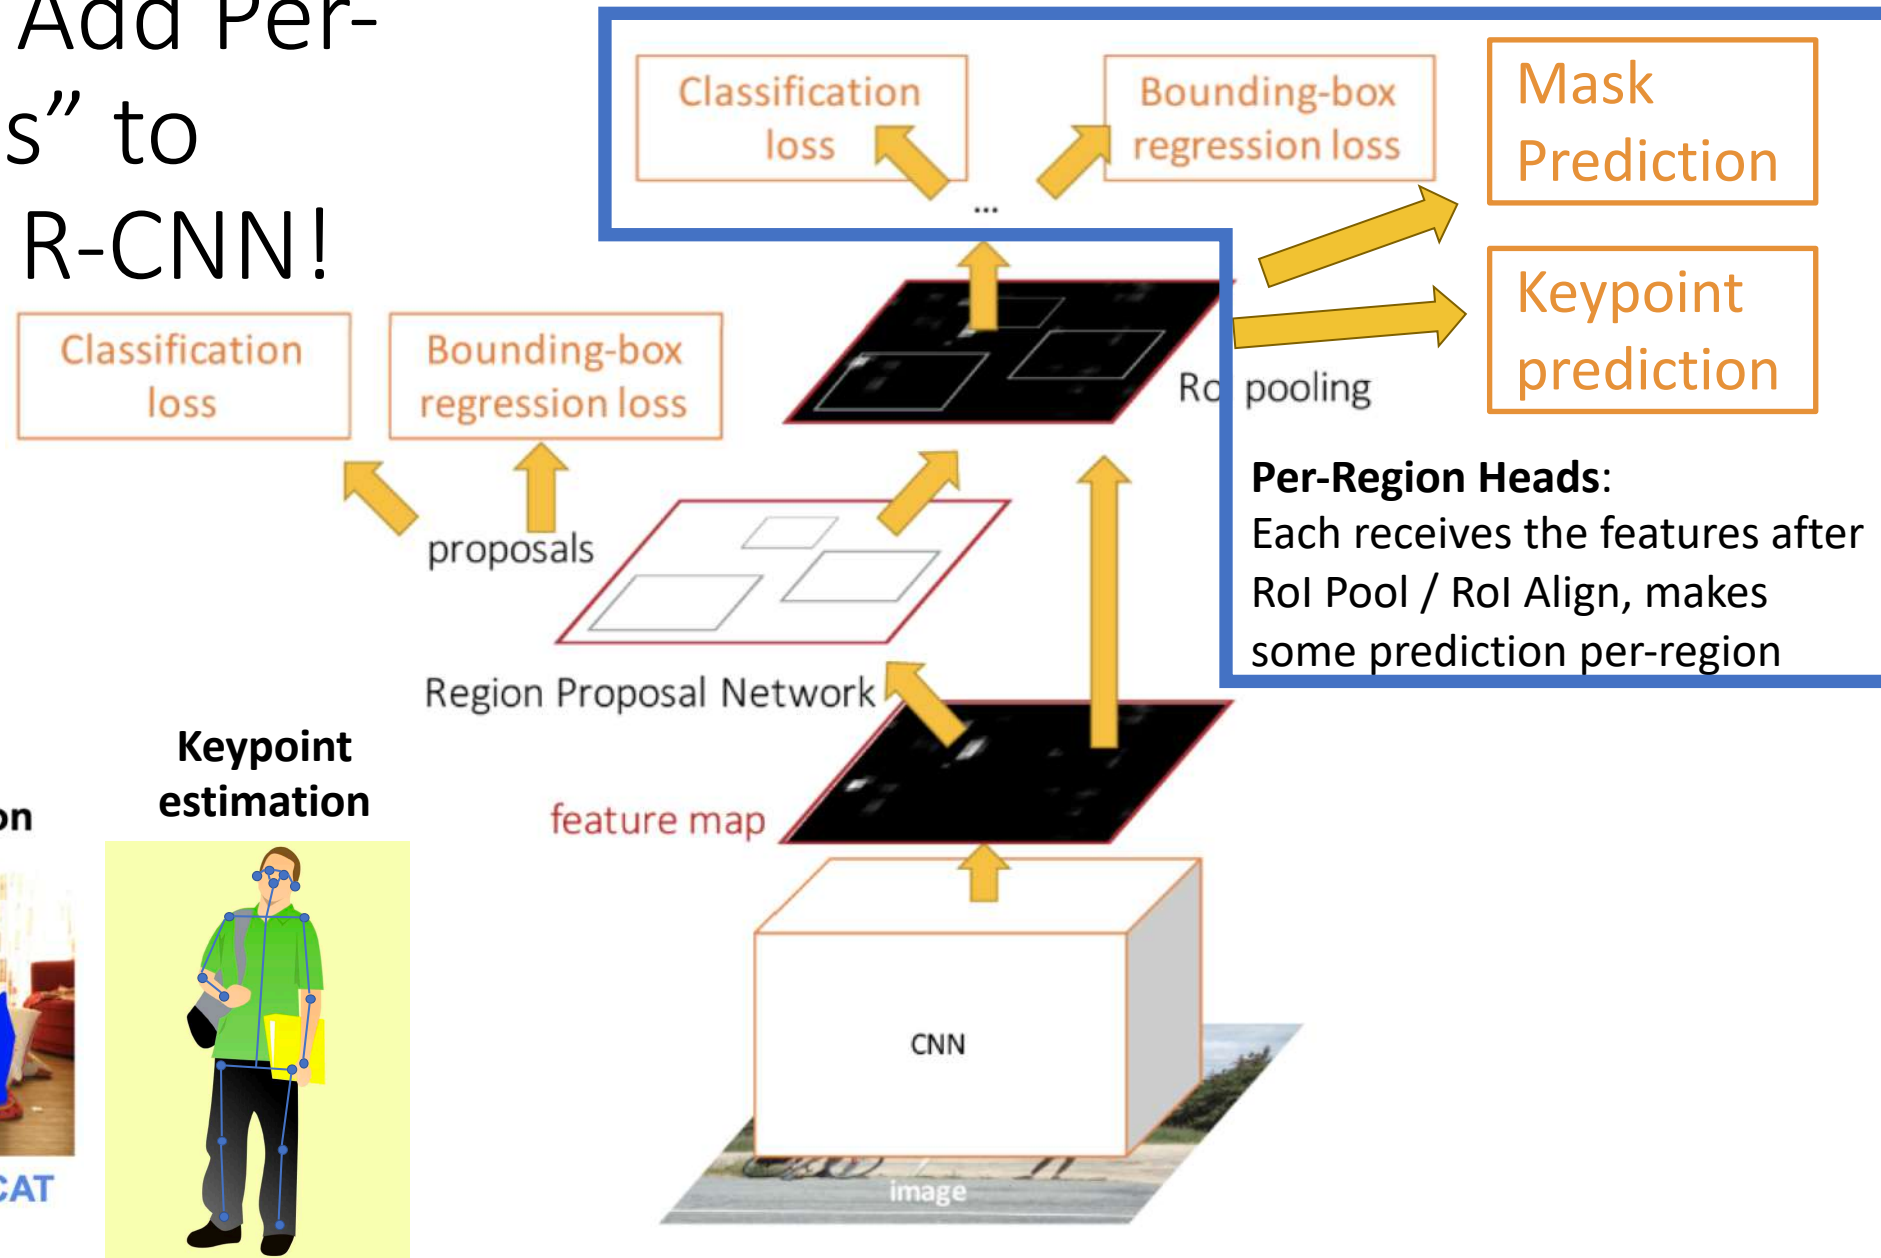


# Joint Instance Segmentation and Pose Estimation



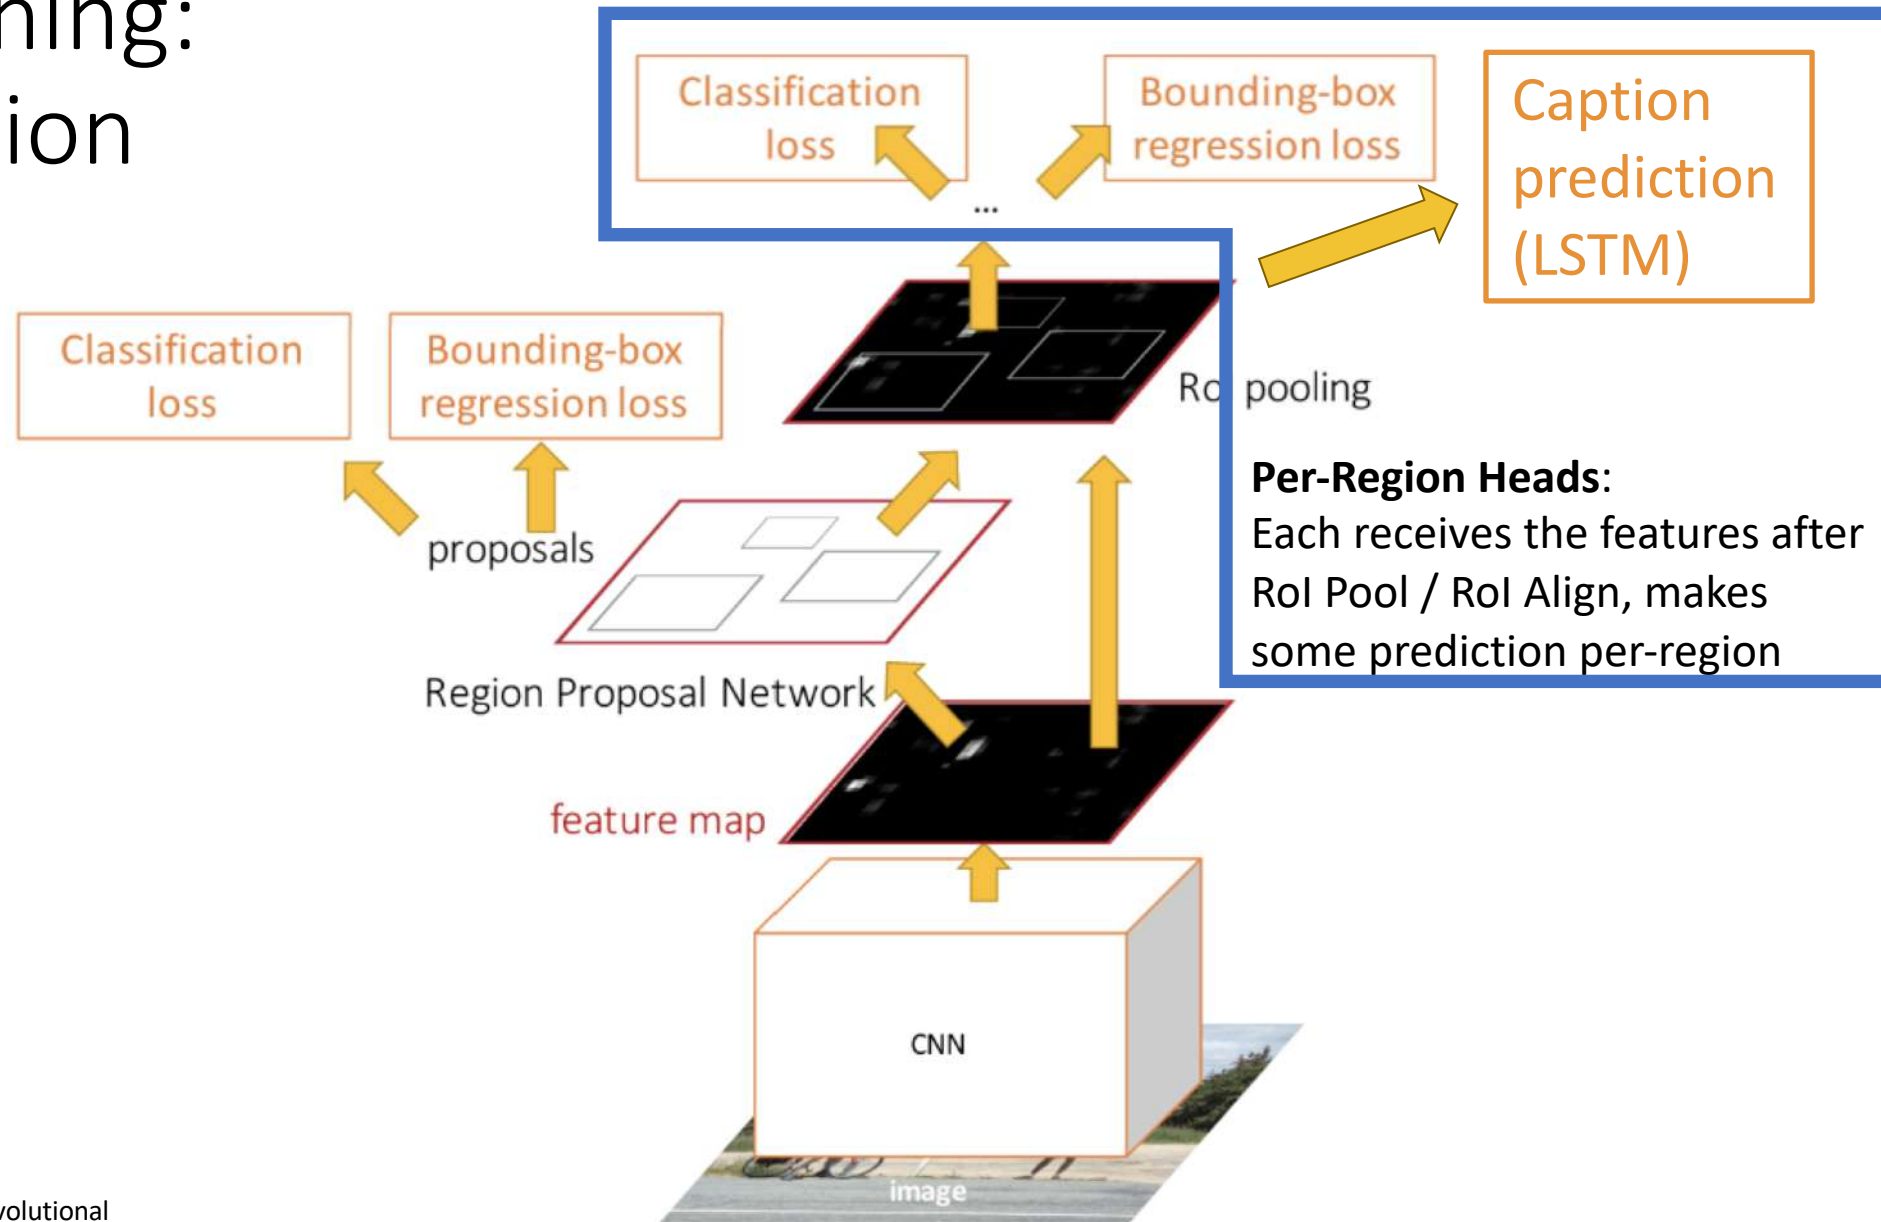
He et al, "Mask R-CNN", ICCV 2017

# General Idea: Add Per-Region “Heads” to Faster / Mask R-CNN!



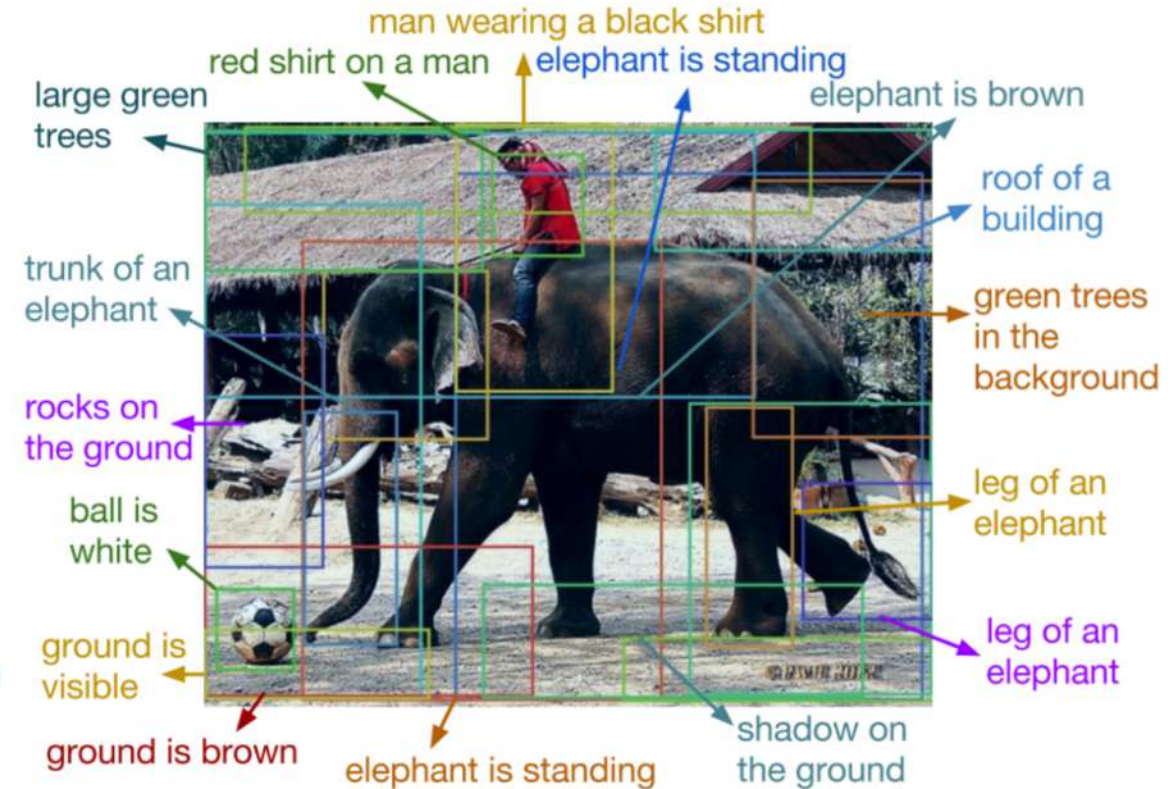
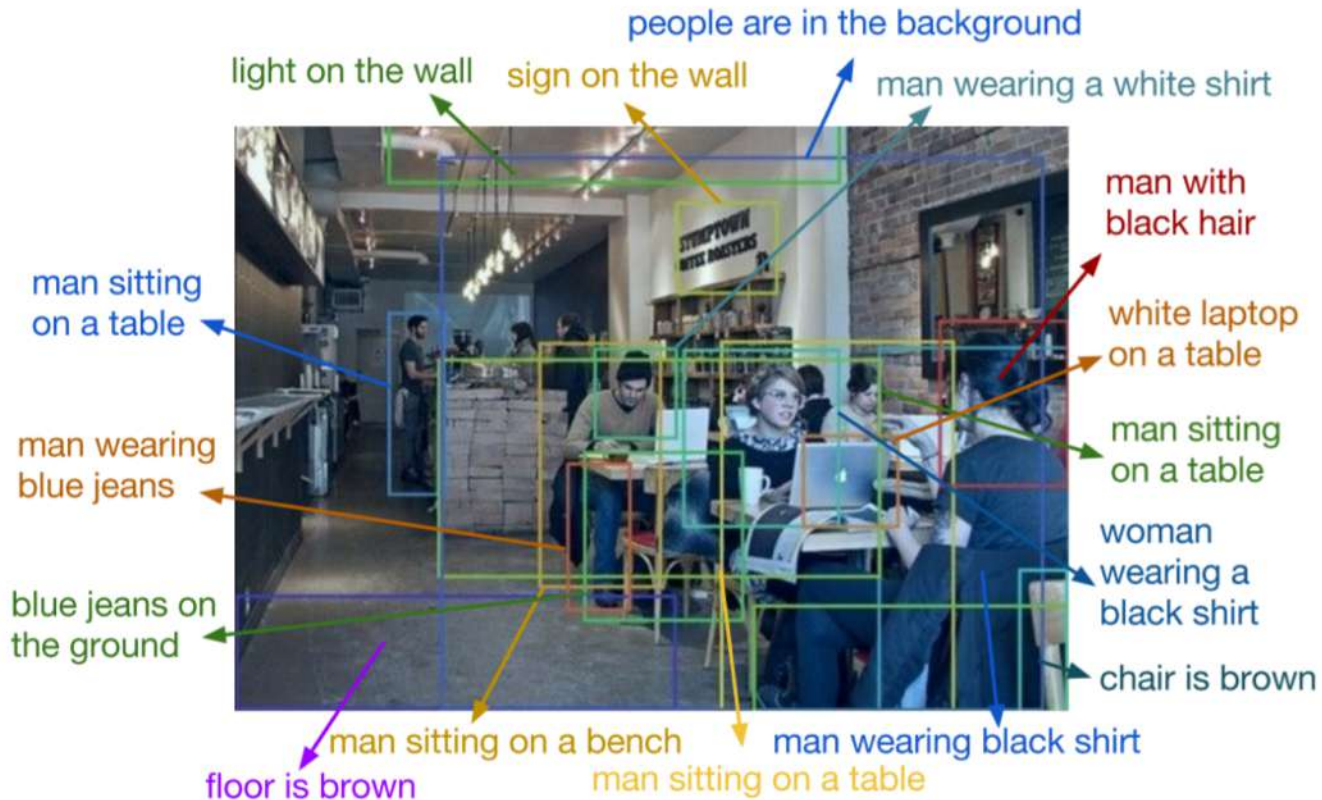


# Dense Captioning: Predict a caption per region!

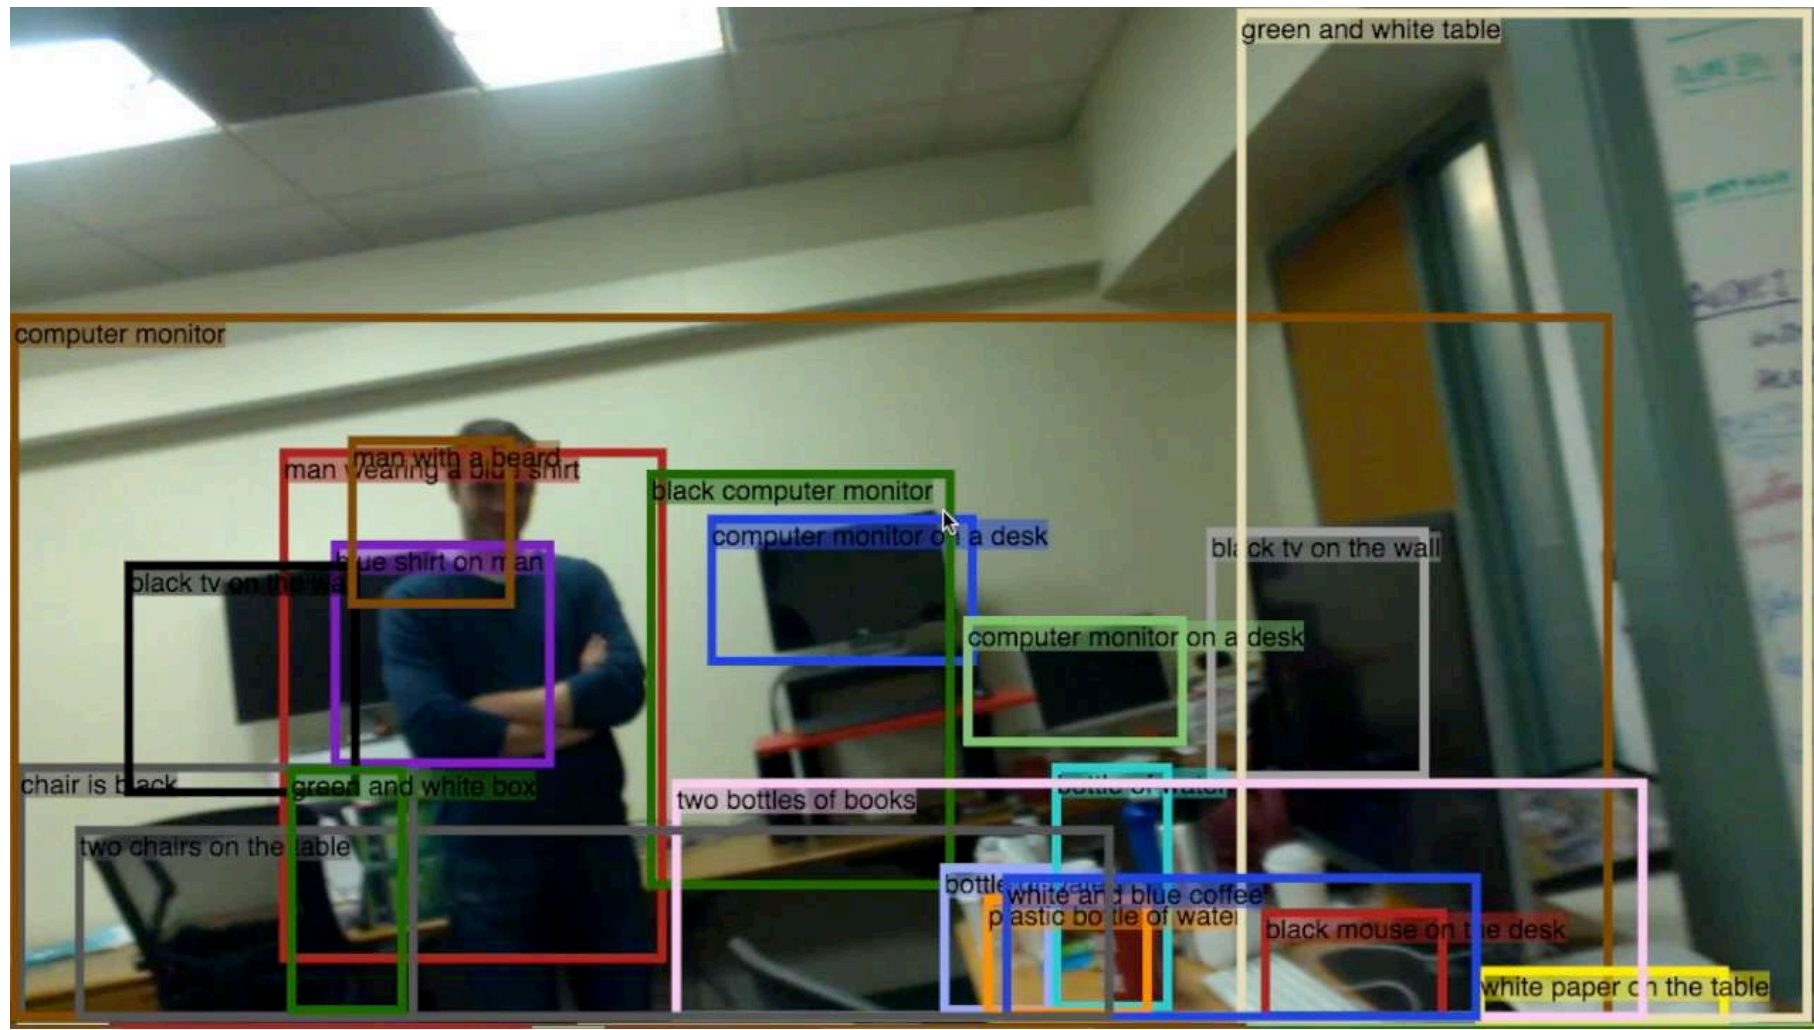


Johnson, Karpathy, and Fei-Fei, "DenseCap: Fully Convolutional Localization Networks for Dense Captioning", CVPR 2016

# Dense Captioning



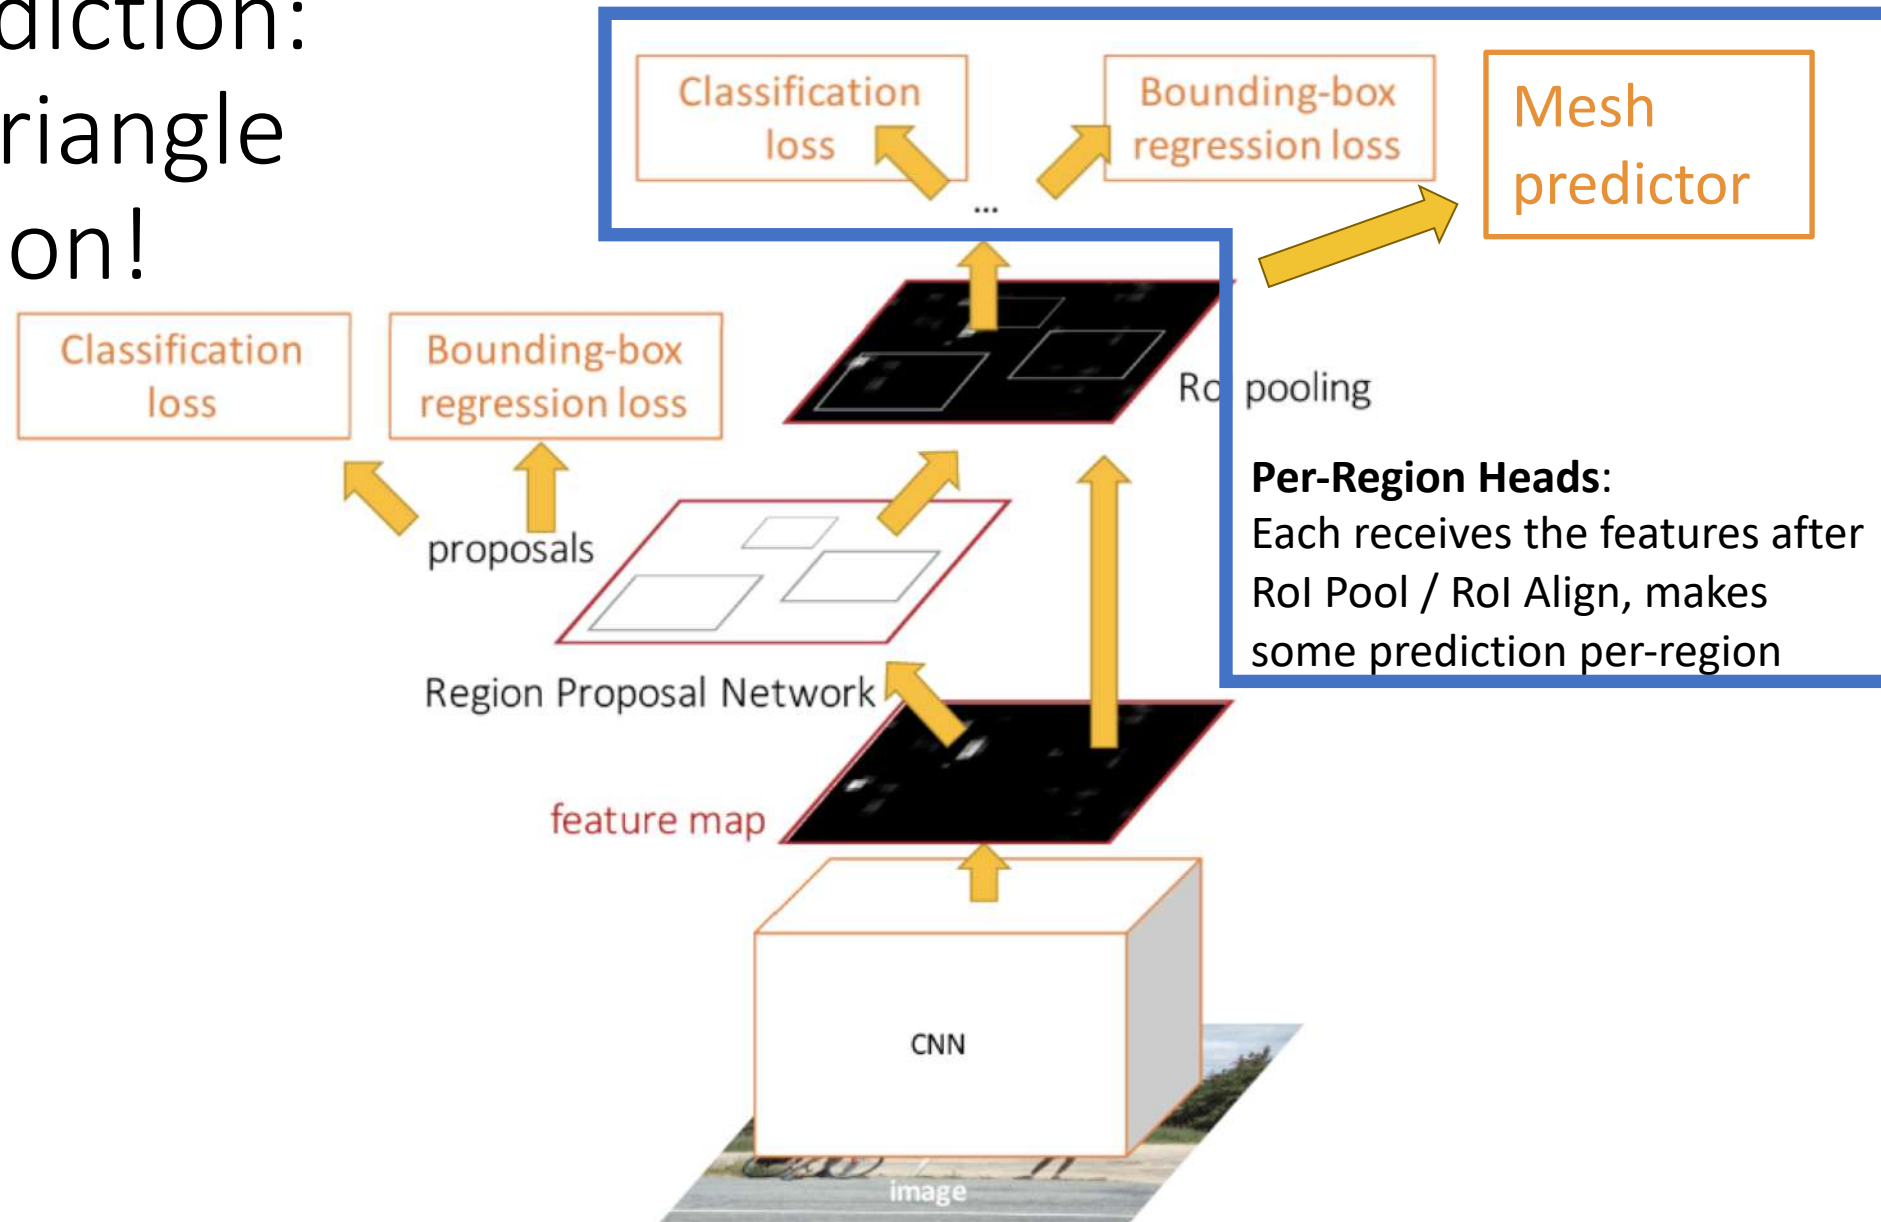
# Dense Captioning



Johnson, Karpathy, and Fei-Fei, "DenseCap: Fully Convolutional Localization Networks for Dense Captioning", CVPR 2016





# 3D Shape Prediction: Predict a 3D triangle mesh per region!





# 3D Shape Prediction: Mask R-CNN + Mesh Head

Mask R-CNN:  
2D Image -> 2D shapes



**Mesh R-CNN:**  
2D Image -> **3D** shapes



More details  
next time!

He, Gkioxari, Dollár, and Girshick, "Mask R-CNN", ICCV 2017

Gkioxari, Malik, and Johnson, "Mesh R-CNN", ICCV 2019



# Summary: Many Computer Vision Tasks!


## Classification



**CAT**

No spatial extent


## Semantic Segmentation



**GRASS, CAT, TREE, SKY**

No objects, just pixels

## Object Detection



**DOG, DOG, CAT**



Multiple Objects

## Instance Segmentation



**DOG, DOG, CAT**

[This image is CC0 public domain](#)



# Next Time: 3D Vision